

dl_project

May 19, 2024

1 The aim of this project is to build a deep neural network capable of analyzing images captured by smartphones to quickly and accurately identify crop diseases.

We are analyzing 54,306 images of plant leaves, which are categorized into 38 class labels. Each class label corresponds to a crop-disease pair, and our goal is to predict this pair from the image of the plant leaf. These images are sourced from the dataset available at the following repository: https://github.com/digitalepidemiologylab/plantvillage_deeplearning_paper_dataset.

To train our AI-based image recognition system, we will utilize this dataset. In all our experiments, we utilize three different versions of the PlantVillage dataset. We begin with the original dataset in color, then we explore a grayscale version, and finally, we conduct our experiments on a version where the leaves are segmented. This approach allows us to assess the performance and robustness of our image recognition system in various contexts. We analyze how variations such as color, grayscale, and leaf segmentation can impact the model's results. By understanding how our system behaves under these different conditions, we can better evaluate its ability to generalize and operate effectively in real-world environments. These three versions of the data are already available via the above-mentioned link.

The different of crop disease types used in this project :

- 0: Grape_ ___ healthy
- 1: Peach_ ___ Bacterial_ spot
- 2: Apple_ ___ healthy
- 3: Orange_ ___ Haunglongbing_ (Citrus_ greening)
- 4: Corn_(maize)_ ___ healthy
- 5: Tomato_ ___ Septoria_ leaf_ spot
- 6: Tomato_ ___ healthy
- 7: Corn_(maize)_ ___ Common_ rust\$
- 8: Tomato_ ___ Early_ blight
- 9: Potato_ ___ Late_ blight
- 10: Peach_ ___ healthy
- 11: Corn_(maize)_ ___ Northern_ Leaf_ Blight

12: Blueberry ____ healthy
13: Grape ____ Leaf_blight_(Isariopsis_Leaf_Spot)
14: Tomato ____ Leaf_Mold
15: Soybean ____ healthy
16: Cherry_(including_sour) ____ healthy
17: Tomato ____ Spider_mites Two-spotted_spider_mite
18: Potato ____ healthy
19: Corn_(maize) ____ Cercospora_leaf_spot_Gray_leaf_spot
20: Cherry_(including_sour) ____ Powdery_mildew
21: Apple ____ Cedar_apple_rust
22: Squash ____ Powdery_mildew
23: Tomato ____ Late_blight
24: Grape ____ Black_rot
25: Pepper,_bell ____ healthy
26: Tomato ____ Target_Spot
27: Apple ____ Black_rot
28: Tomato ____ Bacterial_spot
29: Strawberry ____ healthy
30: Pepper,_bell ____ Bacterial_spot
31: Raspberry ____ healthy
32: Tomato ____ Tomato_Yellow_Leaf_Curl_Virus
33: Apple ____ Apple_scab
34: Potato ____ Early_blight
35: Tomato ____ Tomato_mosaic_virus
36: Strawberry ____ Leaf_scorch
37: Grape ____ Esca_(Black_Measles)

```
[1]: import numpy as np
import pandas as pd
import torch
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
import torch.utils.data as data
import matplotlib.pyplot as plt
```

```

from collections import defaultdict
import seaborn as sns

from PIL import Image, ImageFile
from torchvision import datasets, transforms, utils
from torch.utils.data import ConcatDataset
from torch.utils.data import DataLoader

from sklearn import manifold
from sklearn import decomposition
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay

import subprocess
import shutil

import os
import random
import time

```

#Understanding the data

[2]: SEED = 1

```

random.seed(SEED)
np.random.seed(SEED)
torch.manual_seed(SEED)
torch.cuda.manual_seed(SEED)
torch.backends.cudnn.deterministic = True

```

[3]: ImageFile.LOAD_TRUNCATED_IMAGES = True

```

def extract_folder(repo_url, clone_dir, folder_name):
    if os.path.exists(folder_name):
        os.chdir("../")
        shutil.move(os.path.join(clone_dir, folder_name), folder_name)
        print(f"Folder '{folder_name}' extracted successfully.")

    total_files = 0
    classes = []
    # Count number of files in each directory
    for root, dirs, files in os.walk(folder_name):
        for directory in dirs:
            dir_path = os.path.join(root, directory)
            classes.append(directory)
            num_files = len([name for name in os.listdir(dir_path) if os.
            path.isfile(os.path.join(dir_path, name))])
            total_files += num_files

```

```

        print(f"Directory '{directory}' contains {num_files} files.")

    print(f"Total number of files: {total_files}")

else:
    print(f"Folder '{folder_name}' not found in the repository.")
    classes = []

print()
print(f"total number of classes : '{len(classes)}'.")
print(classes)

return classes

```

```
[ ]: # Clone the repository
repo_url = "https://github.com/digitalepidemiologylab/
    ->plantvillage_deeplearning_paper_dataset.git"
clone_dir = "plantvillage_deeplearning_paper_dataset"
subprocess.run(["git", "clone", repo_url, clone_dir])
os.chdir(clone_dir)

# Raw color images
classes = extract_folder(repo_url, clone_dir, "raw/color")
```

Folder 'raw/color' extracted successfully.

Directory 'Tomato___Early_blight' contains 1000 files.

Directory 'Apple___Black_rot' contains 621 files.

Directory 'Tomato___healthy' contains 1591 files.

Directory 'Peach___healthy' contains 360 files.

Directory 'Strawberry___healthy' contains 456 files.

Directory 'Apple___healthy' contains 1645 files.

Directory 'Tomato___Late_blight' contains 1909 files.

Directory 'Corn_(maize)___healthy' contains 1162 files.

Directory 'Corn_(maize)___Common_rust_' contains 1192 files.

Directory 'Pepper,_bell___Bacterial_spot' contains 997 files.

Directory 'Soybean___healthy' contains 5090 files.

Directory 'Grape___Leaf_blight_(Isariopsis_Leaf_Spot)' contains 1076 files.

Directory 'Strawberry___Leaf_scorch' contains 1109 files.

Directory 'Tomato___Tomato_Yellow_Leaf_Curl_Virus' contains 5357 files.

Directory 'Potato___healthy' contains 152 files.

Directory 'Tomato___Tomato_mosaic_virus' contains 373 files.

Directory 'Orange___Haunglongbing_(Citrus_greening)' contains 5507 files.

Directory 'Corn_(maize)___Northern_Leaf_Blight' contains 985 files.

Directory 'Pepper,_bell___healthy' contains 1478 files.

Directory 'Tomato___Septoria_leaf_spot' contains 1771 files.

Directory 'Grape___Esca_(Black_Measles)' contains 1383 files.

Directory 'Tomato___Target_Spot' contains 1404 files.

Directory 'Tomato___Leaf_Mold' contains 952 files.

```

Directory 'Corn_(maize)___Cercospora_leaf_spot_Gray_leaf_spot' contains 513
files.
Directory 'Cherry_(including_sour)___Powdery_mildew' contains 1052 files.
Directory 'Apple___Cedar_apple_rust' contains 275 files.
Directory 'Raspberry___healthy' contains 371 files.
Directory 'Grape___healthy' contains 423 files.
Directory 'Apple___Apple_scab' contains 630 files.
Directory 'Peach___Bacterial_spot' contains 2297 files.
Directory 'Cherry_(including_sour)___healthy' contains 854 files.
Directory 'Tomato___Bacterial_spot' contains 2127 files.
Directory 'Grape___Black_rot' contains 1180 files.
Directory 'Squash___Powdery_mildew' contains 1835 files.
Directory 'Potato___Early_blight' contains 1000 files.
Directory 'Blueberry___healthy' contains 1502 files.
Directory 'Potato___Late_blight' contains 1000 files.
Directory 'Tomato___Spider_mites Two-spotted_spider_mite' contains 1676 files.
Total number of files: 54305

```

```

total number of classes : '38'.
['Tomato___Early_blight', 'Apple___Black_rot', 'Tomato___healthy',
'Peach___healthy', 'Strawberry___healthy', 'Apple___healthy',
'Tomato___Late_blight', 'Corn_(maize)___healthy', 'Corn_(maize)___Common_rust_',
'Pepper,_bell___Bacterial_spot', 'Soybean___healthy',
'Grape___Leaf_blight_(Isariopsis_Leaf_Spot)', 'Strawberry___Leaf_scorch',
'Tomato___Tomato_Yellow_Leaf_Curl_Virus', 'Potato___healthy',
'Tomato___Tomato_mosaic_virus', 'Orange___Haunglongbing_(Citrus_greening)',
'Corn_(maize)___Northern_Leaf_Blight', 'Pepper,_bell___healthy',
'Tomato___Septoria_leaf_spot', 'Grape___Esca_(Black_Measles)',
'Tomato___Target_Spot', 'Tomato___Leaf_Mold',
'Corn_(maize)___Cercospora_leaf_spot_Gray_leaf_spot',
'Cherry_(including_sour)___Powdery_mildew', 'Apple___Cedar_apple_rust',
'Raspberry___healthy', 'Grape___healthy', 'Apple___Apple_scab',
'Peach___Bacterial_spot', 'Cherry_(including_sour)___healthy',
'Tomato___Bacterial_spot', 'Grape___Black_rot', 'Squash___Powdery_mildew',
'Potato___Early_blight', 'Blueberry___healthy', 'Potato___Late_blight',
'Tomato___Spider_mites Two-spotted_spider_mite']

```

Unbalanced Distribution: The distribution appears to be highly unbalanced. This conclusion is drawn from the significant disparity in the heights of the bars representing different crop disease classes. Some classes have a much larger number of files compared to others.

Imbalance Severity: Classes such as "Tomato_Tomato_Yellow_Leaf_Curl_Virus", "Orange_Haunglongbing_(Citrus_greening)", and "Soybean_healthy" have notably higher numbers of files compared to the rest. Conversely, classes like "Potato_healthy", "Apple_Cedar_apple_rust", and "Corn_(maize)_Cercospora_leaf_spot_Gray_leaf_spot" have relatively fewer files.

Impact on Model Performance: The presence of highly unbalanced classes can potentially impact the performance of machine learning models trained on this dataset. Models may have a tendency

to be biased towards the majority classes, leading to poorer performance on minority classes.

Data Augmentation or Sampling Strategies: Addressing class imbalance may require techniques such as data augmentation for minority classes.

That leads to improved model performance and reduced overfitting. However, it's crucial to select augmentation techniques carefully and evaluate models appropriately to avoid introducing bias

#A. Data Preprocessing.

```
[6]: class CropDiseaseDataset(data.Dataset):
    def __init__(self, root_dir, train=True, validation=False, gray_scale=False, ↴segmented=False):
        """Initializes a dataset containing images and labels."""
        super().__init__()
        self.gray_scale = gray_scale
        self.root_dir = root_dir
        self.train = train
        self.validation = validation
        self.segmented = segmented

        # Get all jpg images
        self.files = []
        for class_name in os.listdir(self.root_dir):
            class_dir = os.path.join(self.root_dir, class_name)
            # Filter files based on the segmented condition
            if self.segmented:
                files = [os.path.join(class_dir, file) for file in os.listdir(class_dir) if file.endswith(".jpg")]
            else:
                files = [os.path.join(class_dir, file) for file in os.listdir(class_dir) if file.endswith(".JPG")]
            self.files.extend(files)

        random.shuffle(self.files)

        # Define the split ratios
        train_ratio = 0.65
        validation_ratio = 0.15

        self.classes = []
        for root, dirs, _ in os.walk(root_dir):
            for directory in dirs:
                self.classes.append(directory)

        # Split the dataset into train, validation, and test sets
        if self.train:
            num_train_files = int(len(self.files) * train_ratio)
            self.files = self.files[:num_train_files]
```

```

        elif self.validation:
            num_train_files = int(len(self.files) * train_ratio)
            num_validation_files = int(len(self.files) * validation_ratio)
            self.files = self.files[num_train_files:num_train_files + num_validation_files]
        else:
            num_train_files = int(len(self.files) * train_ratio)
            num_validation_files = int(len(self.files) * validation_ratio)
            self.files = self.files[num_train_files + num_validation_files:]

    # Remove all invalid files
    def is_valid(filename):
        """Check that a file is not corrupted, is in RGB, and has non-zero size"""
        if os.path.getsize(filename) == 0:
            return False
        valid_extensions = ['.jpg']
        _, ext = os.path.splitext(filename)
        if ext.lower() not in valid_extensions:
            return False
        try:
            with Image.open(filename, 'r') as img:
                img.verify()
                mode = img.mode
                return mode == 'RGB'
        except:
            return False

    wrong = []
    for i in range(len(self.files)):
        if not is_valid(self.files[i]):
            wrong.append(i)
    self.files = np.delete(self.files, wrong)
    random.shuffle(self.files)

    if self.gray_scale:
        self.means, self.stds = self.calculate_mean_std_gray()
        self.means = np.array(self.means)
        self.stds = np.array(self.stds)
    else:
        self.means, self.stds = self.calculate_mean_std()

    self.transform = self.get_transform()

def calculate_mean_std(self):
    sum_channel = np.zeros(3)
    squared_sum_channel = np.zeros(3)

```

```

total_pixels = 0
total_images = len(self.files)

# Iterate through all images to accumulate sums
for file in self.files:
    img = Image.open(file)

    img_array = np.array(img)
    total_pixels += img_array.size / 3
    sum_channel += np.sum(img_array, axis=(0, 1)) / 255.0
    squared_sum_channel += np.sum((img_array / 255.0) ** 2, axis=(0, 1))

means = sum_channel / total_pixels
stds = np.sqrt((squared_sum_channel / total_pixels) - (means ** 2))

return means.tolist(), stds.tolist()

def calculate_mean_std_gray(self):
    sum_channel = np.zeros(1) # Grayscale has only one channel
    squared_sum_channel = np.zeros(1)
    total_pixels = 0
    total_images = len(self.files)

    # Iterate through all images to accumulate sums
    for file in self.files:
        img = Image.open(file).convert("L") # Convert image to grayscale

        img_array = np.array(img)
        total_pixels += img_array.size
        sum_channel += np.sum(img_array) / 255.0
        squared_sum_channel += np.sum((img_array / 255.0) ** 2)

    means = sum_channel / total_pixels
    stds = np.sqrt((squared_sum_channel / total_pixels) - (means ** 2))

    return [means], [stds]

def get_transform(self):
    if self.gray_scale:
        return transforms.Compose([
            transforms.RandomVerticalFlip(),
            transforms.RandomHorizontalFlip(),
            transforms.RandomApply([
                transforms.ColorJitter(
                    brightness=torch.rand(1).item(),
                    contrast=torch.rand(1).item(),
                    saturation=torch.rand(1).item(),

```

```

                hue=torch.rand(1).item() * 0.5,
            )
        ]),
        transforms.Resize((120, 120)),
        transforms.Grayscale(num_output_channels=1), # Convert to u
↳grayscale
        transforms.ToTensor(),
        transforms.Normalize(self.means, self.stds)
    ))
else:
    return transforms.Compose([
        transforms.RandomVerticalFlip(),
        transforms.RandomHorizontalFlip(),
        transforms.RandomApply([
            transforms.ColorJitter(
                brightness=torch.rand(1).item(),
                contrast=torch.rand(1).item(),
                saturation=torch.rand(1).item(),
                hue=torch.rand(1).item() * 0.5,
            )
        ]),
        transforms.Resize((120, 120)),
        transforms.ToTensor(),
        transforms.Normalize(self.means, self.stds)
    ))
}

def __len__(self):
    """Returns the size of the dataset."""
    return len(self.files)

def __getitem__(self, index):
    """Returns the index-th data item of the dataset."""
    if index < 0 or index >= self.__len__():
        raise IndexError(
            'Wrong index (must be in [{}], {}, given: {})'
            .format(0, self.__len__(), index)
        )

    # Load image and transform
    img = Image.open(self.files[index])
    img = self.transform(img)

    # Get label
    class_name = os.path.basename(os.path.dirname(self.files[index]))
    label = self.classes.index(class_name)

    return img, torch.tensor(label)

```

```
[ ]: # Provided data: directory names and corresponding file counts
data_ = {
    'Corn_(maize)___healthy': 1162,
    'Grape___Esca_(Black_Measles)': 1383,
    'Orange___Haunglongbing_(Citrus_greening)': 5507,
    'Tomato___healthy': 1591,
    'Tomato___Bacterial_spot': 2127,
    'Pepper,_bell___Bacterial_spot': 997,
    'Apple___Black_rot': 621,
    'Potato___Early_blight': 1000,
    'Tomato___Late_blight': 1909,
    'Cherry_(including_sour)___healthy': 854,
    'Tomato___Septoria_leaf_spot': 1771,
    'Grape___Leaf_blight_(Isariopsis_Leaf_Spot)': 1076,
    'Cherry_(including_sour)___Powdery_mildew': 1052,
    'Potato___healthy': 152,
    'Strawberry___Leaf_scorch': 1109,
    'Grape___healthy': 423,
    'Peach___healthy': 360,
    'Tomato___Spider_mites Two-spotted_spider_mite': 1676,
    'Apple___Apple_scab': 630,
    'Raspberry___healthy': 371,
    'Tomato___Early_blight': 1000,
    'Tomato___Target_Spot': 1404,
    'Apple___Cedar_apple_rust': 275,
    'Corn_(maize)___Northern_Leaf_Blight': 985,
    'Soybean___healthy': 5090,
    'Strawberry___healthy': 456,
    'Tomato___Tomato_Yellow_Leaf_Curl_Virus': 5357,
    'Peach___Bacterial_spot': 2297,
    'Grape___Black_rot': 1180,
    'Apple___healthy': 1645,
    'Tomato___Tomato_mosaic_virus': 373,
    'Blueberry___healthy': 1502,
    'Tomato___Leaf_Mold': 952,
    'Corn_(maize)___Common_rust_': 1192,
    'Squash___Powdery_mildew': 1835,
    'Pepper,_bell___healthy': 1478,
    'Potato___Late_blight': 1000,
    'Corn_(maize)___Cercospora_leaf_spot_Gray_leaf_spot': 513
}

directories = list(data_.keys())
file_counts = list(data_.values())
colors = sns.color_palette("husl", len(directories))

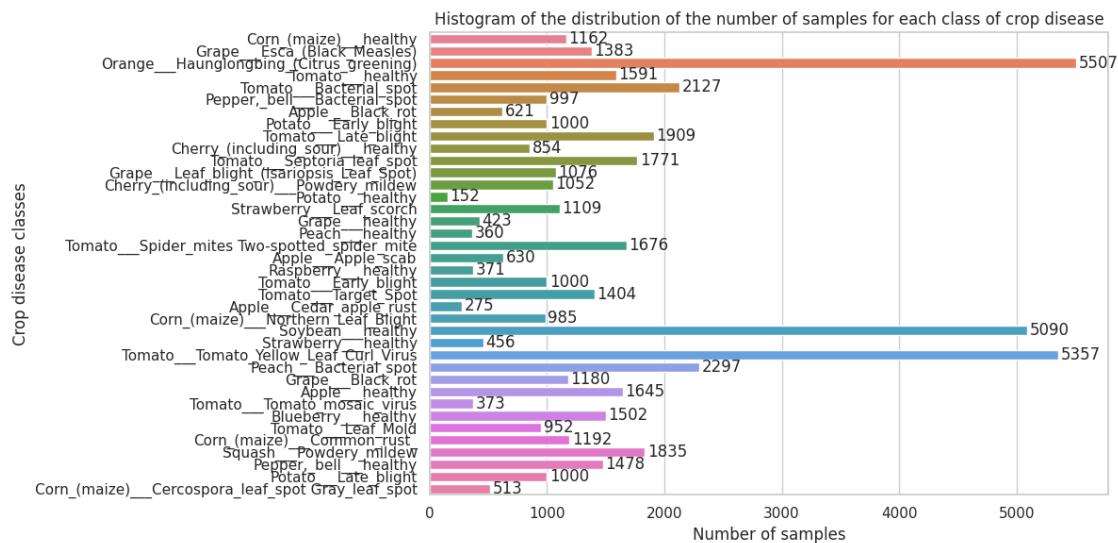
sns.set(style="whitegrid")
```

```

plt.figure(figsize=(12, 6))
sns.barplot(x=file_counts, y=directories, hue=directories, dodge=False, palette=colors, legend=False)
for i, count in enumerate(file_counts):
    plt.text(count + 20, i, str(count), va='center')

plt.xlabel('Number of samples')
plt.ylabel('Crop disease classes')
plt.title('Histogram of the distribution of the number of samples for each class of crop disease')
plt.tight_layout()
plt.show()

```

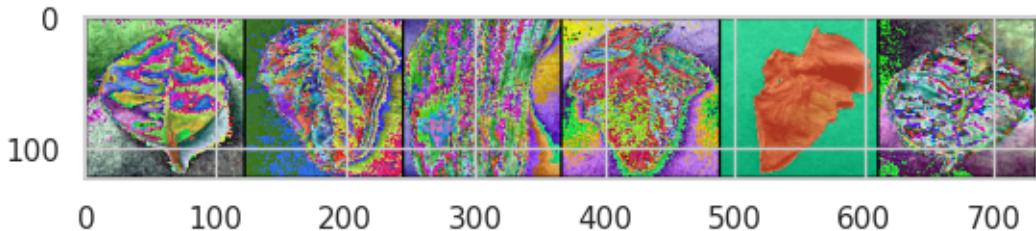


```
[ ]: def show_images(img):
    plt.imshow(transforms.functional.to_pil_image(img))
    plt.show()
```

Checking out.

```
[ ]: extracted_folder = "raw/color"
my_dataset = CropDiseaseDataset(root_dir=extracted_folder, train=True, gray_scale=False, segmented = False)
train_loader = DataLoader(my_dataset, batch_size=6, shuffle=True, num_workers=2)

images, labels = next(iter(train_loader))
show_images(utils.make_grid(images))
print(*[classes[l] for l in labels])
```



Orange___Haunglongbing_(Citrus_greening) Tomato___Bacterial_spot
 Corn_(maize)___Northern_Leaf_Blight Tomato___Bacterial_spot Tomato___Late_blight
 Cherry_(including_sour)___Powdery_mildew

```
[ ]: my_dataset = CropDiseaseDataset(root_dir=extracted_folder, train=True, ↴  

    ↪gray_scale=True, segmented=False)  

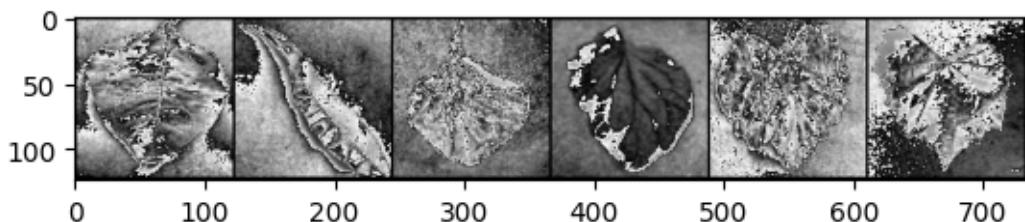
train_loader = DataLoader(my_dataset, batch_size=6, shuffle=True, num_workers=2)  
  

images, labels = next(iter(train_loader))  

show_images(utils.make_grid(images))  

print(*[my_dataset.classes[l] for l in labels])
```

Tomato___Tomato_Yellow_Leaf_Curl_Virus Soybean___healthy
 Grape___Esca_(Black_Measles) Orange___Haunglongbing_(Citrus_greening)
 Apple___healthy Pepper,_bell___healthy



Pepper,_bell___healthy Peach___Bacterial_spot Soybean___healthy
 Tomato___Tomato_Yellow_Leaf_Curl_Virus Grape___Black_rot Grape___Black_rot

#B. Defining the convolutional network architecture.

Initialize our parameters, determine the number of model parameters and a function to tell us how long an epoch takes.

```
[7]: def initialize_parameters(m):  

    if isinstance(m, nn.Conv2d):  

        nn.init.kaiming_normal_(m.weight.data, nonlinearity='relu')  

        nn.init.constant_(m.bias.data, 0)
```

```

    elif isinstance(m, nn.Linear):
        nn.init.xavier_normal_(m.weight.data, gain=nn.init.
→calculate_gain('relu'))
        nn.init.constant_(m.bias.data, 0)

def count_parameters(model):
    return sum(p.numel() for p in model.parameters() if p.requires_grad)

def epoch_time(start_time, end_time):
    elapsed_time = end_time - start_time
    elapsed_mins = int(elapsed_time / 60)
    elapsed_secs = int(elapsed_time - (elapsed_mins * 60))
    return elapsed_mins, elapsed_secs

```

```

[8]: class CNN(nn.Module):
    def __init__(self, grayscale=False):
        super(CNN, self).__init__()
        if grayscale:
            in_channels = 1 # Set input channels to 1 for grayscale images
        else:
            in_channels = 3 # Keep input channels as 3 for RGB images
        self.conv1 = nn.Conv2d(in_channels, 32, kernel_size=3, stride=1, □
→padding=1)
        self.conv2 = nn.Conv2d(32, 64, kernel_size=3, stride=1, padding=1)
        self.conv3 = nn.Conv2d(64, 128, kernel_size=3, stride=1, padding=1)
        self.conv4 = nn.Conv2d(128, 256, kernel_size=3, stride=1, padding=1)
        self.pool = nn.MaxPool2d(2)
        self.fc1 = nn.Linear(256 * 7 * 7, 512)
        self.fc2 = nn.Linear(512, 64)
        self.fc3 = nn.Linear(64, 38)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = self.pool(F.relu(self.conv3(x)))
        x = self.pool(F.relu(self.conv4(x)))
        x = x.view(-1, 256 * 7 * 7)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

```

```

[ ]: model = CNN()
print(model)

CNN(
  (conv1): Conv2d(3, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (conv2): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

```

```

(conv3): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(conv4): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
(fc1): Linear(in_features=12544, out_features=512, bias=True)
(fc2): Linear(in_features=512, out_features=64, bias=True)
(fc3): Linear(in_features=64, out_features=38, bias=True)
)

```

2 C. Optimization stage (hyperparameter selections)

```
[ ]: # Splitting the dataset
extracted_folder = "raw/color"
train_set = CropDiseaseDataset(root_dir=extracted_folder, train=True,
                                validation=False, gray_scale=False, segmented=False)
validation_set = CropDiseaseDataset(root_dir=extracted_folder, train=False,
                                    validation=True, gray_scale=False, segmented=False)
test_set = CropDiseaseDataset(root_dir=extracted_folder, train=False,
                               validation=False, gray_scale=False, segmented=False)

trainloader = torch.utils.data.DataLoader(train_set, batch_size=64,
                                         shuffle=True, num_workers=2)
validationloader = torch.utils.data.DataLoader(validation_set, batch_size=64,
                                                shuffle=False, num_workers=2)
testloader = torch.utils.data.DataLoader(test_set, batch_size=64, shuffle=False,
                                         num_workers=2)
```

```
[ ]: print(f"Train set: '{len(train_set)}' images, f"Validation set:{len(validation_set)}' images, f"Test set: '{len(test_set)}' images")
```

Train set: '34321' images, Validation set: '7920' images, Test set: '10562' images

```
[ ]: merged_dataset = ConcatDataset([train_set, validation_set, test_set])

# Create a single loader for the combined dataset
mergedloader = DataLoader(merged_dataset, batch_size=64, shuffle=True,
                           num_workers=2)
```

```
[ ]: class_counts = defaultdict(int)
for images, labels in mergedloader:
    for label in labels:
        class_counts[label.item()] += 1

class_labels = sorted(class_counts.keys())
counts = [class_counts[label] for label in class_labels]
colors = sns.color_palette("husl", len(class_labels))
```

```

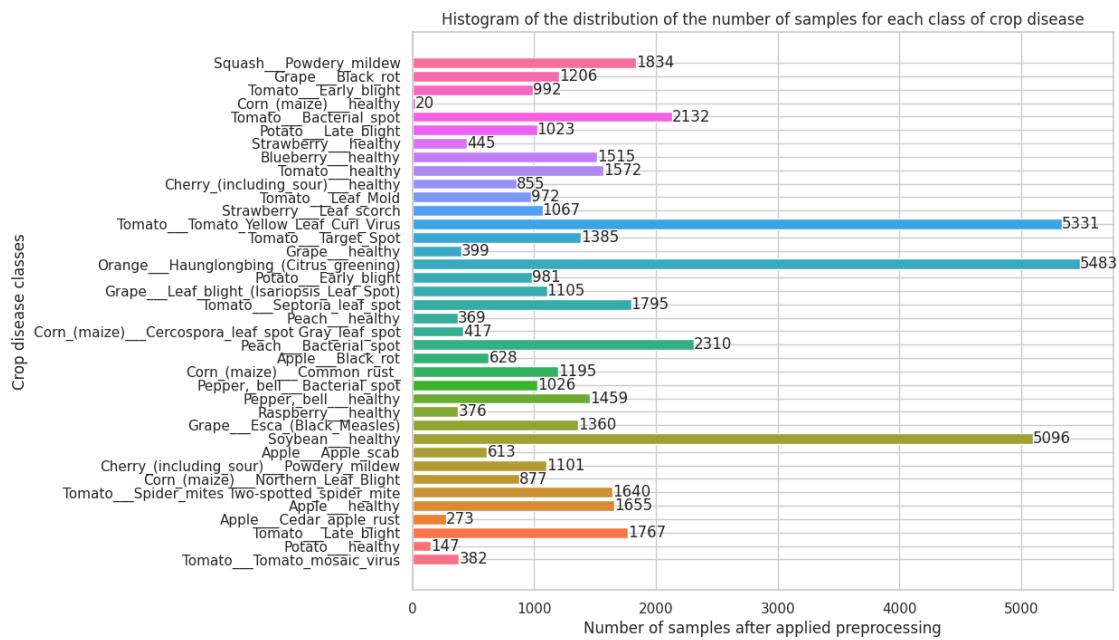
class_names = [train_set.classes[label] for label in class_labels]

sns.set(style="whitegrid")
plt.figure(figsize=(10, 8))
plt.barh(class_names, counts, color=colors)
plt.xlabel('Number of samples after applied preprocessing')
plt.ylabel('Crop disease classes')
plt.title('Histogram of the distribution of the number of samples for each class of crop disease')

for i, (count, name) in enumerate(zip(counts, class_names)):
    plt.text(count, i, str(count), ha='left', va='center')

plt.show()

```



```

[9]: device = 'cpu'
if torch.cuda.is_available():
    device = 'cuda'

# Define the training method
def train(model, num_epochs, trainloader, testloader, criterion, optimizer,
validation_phase=False):
    train_avg_loss = []
    test_avg_loss = []
    test_accuracy = []

```

```

train_accuracy = []

for i in range(num_epochs):
    train_losses = []
    test_losses = []
    correct_train = 0
    total_train = 0

    for x, y in trainloader:
        start_time = time.monotonic()
        x, y = x.to(device), y.to(device)

        pred = model(x)
        loss = criterion(pred, y)
        train_losses.append(loss.detach())

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        # Compute training accuracy
        y_pred_train = pred.argmax(dim=-1)
        correct_train += (y_pred_train == y).sum().item()
        total_train += y.size(0)

    with torch.no_grad():
        correct_test = 0
        total_test = 0

        # Choose loader based on phase
        if validation_phase:
            loader = validationloader
        else:
            loader = testloader

        for x, y in loader:
            x, y = x.to(device), y.to(device)

            pred = model(x)
            loss = criterion(pred, y)
            test_losses.append(loss)

            y_pred_test = pred.argmax(dim=-1)
            correct_test += (y_pred_test == y).sum().item()
            total_test += y.size(0)

        # Compute accuracy

```

```

accuracy_train = correct_train / total_train
accuracy_test = correct_test / total_test
test_accuracy.append(accuracy_test)

train_avg_loss.append(torch.mean(torch.tensor(train_losses)).detach())
test_avg_loss.append(torch.mean(torch.tensor(test_losses)).detach())

if validation_phase:
    print("Epoch [{}/{}], Validation Loss: {:.4f}, Validation Accuracy: {:.2f}%, Train Loss: {:.4f}, Train Accuracy: {:.2f}%".format(
        i+1, num_epochs, test_avg_loss[-1], accuracy_test*100,
        train_avg_loss[-1], accuracy_train*100))
else:
    print("Epoch [{}/{}], Test Loss: {:.4f}, Test Accuracy: {:.2f}%, Train Loss: {:.4f}, Train Accuracy: {:.2f}%".format(
        i+1, num_epochs, test_avg_loss[-1], accuracy_test*100,
        train_avg_loss[-1], accuracy_train*100))

end_time = time.monotonic()
epoch_mins, epoch_secs = epoch_time(start_time, end_time)

print(f'| Epoch Time: {epoch_mins}m {epoch_secs}s')

train_accuracy.append(accuracy_train)

# Save the model at the end of the training
torch.save(model.state_dict(), 'model.pt')

return train_avg_loss, test_avg_loss, test_accuracy, train_accuracy

```

```
[10]: # define a method for plotting the results
def plot_training_results(train_avg_loss, validation_avg_loss, validation_accuracy, is_validation=True):
    sns.set(style="whitegrid")
    fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(15, 5))

    if is_validation:
        loss_label = 'Validation Loss'
        accuracy_label = 'Validation Accuracy'
    else:
        loss_label = 'Test Loss'
        accuracy_label = 'Test Accuracy'

    # Plot train and validation/test loss
    ax1.plot(train_avg_loss, label='Train Loss')
    ax1.plot(validation_avg_loss, label=loss_label)
    ax1.set_xlabel('Epoch')
```

```

ax1.set_ylabel('Loss')
ax1.legend(loc='upper right')

# Plot validation/test accuracy
ax2.plot(validation_accuracy, label=accuracy_label)
ax2.set_xlabel('Epoch')
ax2.set_ylabel('Accuracy')
ax2.legend(loc='upper left')

plt.show()

```

Learning rate selection

```

[ ]: num_epochs = 20
learning_rates = [1e-3, 1e-2, 1e-4]

# Lists to store results for each learning rate
train_losses_list = []
validation_losses_list = []
train_accuracy_list = []
validation_accuracy_list = []

for idx, lr in enumerate(learning_rates):
    network = CNN().to(device)
    criterion = nn.CrossEntropyLoss()
    optimizer = torch.optim.AdamW(network.parameters(), lr=lr)

    print(f"Training model for configuration {idx+ 1}...")
    # Train the model
    train_avg_loss, validation_avg_loss, validation_accuracy, train_accuracy = \
        train(network, num_epochs, trainloader, validationloader, criterion, \
        optimizer, validation_phase=True)

    # Store results for plotting
    train_losses_list.append(train_avg_loss)
    validation_losses_list.append(validation_avg_loss)
    train_accuracy_list.append(train_accuracy)
    validation_accuracy_list.append(validation_accuracy)

```

Training model for configuration 1...
Epoch [1/20], Validation Loss: 1.1557, Validation Accuracy: 65.19%, Train Loss: 1.9920, Train Accuracy: 45.28%
| Epoch Time: 0m 36s
Epoch [2/20], Validation Loss: 0.6847, Validation Accuracy: 77.85%, Train Loss: 0.9250, Train Accuracy: 71.22%
| Epoch Time: 0m 36s
Epoch [3/20], Validation Loss: 0.4793, Validation Accuracy: 84.39%, Train Loss: 0.6446, Train Accuracy: 79.49%

```
| Epoch Time: 0m 35s
Epoch [4/20], Validation Loss: 0.4588, Validation Accuracy: 85.06%, Train Loss: 0.5095, Train Accuracy: 83.55%
| Epoch Time: 0m 35s
Epoch [5/20], Validation Loss: 0.3131, Validation Accuracy: 89.44%, Train Loss: 0.4359, Train Accuracy: 85.84%
| Epoch Time: 0m 36s
Epoch [6/20], Validation Loss: 0.3292, Validation Accuracy: 89.00%, Train Loss: 0.3662, Train Accuracy: 87.85%
| Epoch Time: 0m 37s
Epoch [7/20], Validation Loss: 0.3100, Validation Accuracy: 89.90%, Train Loss: 0.3319, Train Accuracy: 89.07%
| Epoch Time: 0m 36s
Epoch [8/20], Validation Loss: 0.3023, Validation Accuracy: 90.03%, Train Loss: 0.3054, Train Accuracy: 89.90%
| Epoch Time: 0m 36s
Epoch [9/20], Validation Loss: 0.2389, Validation Accuracy: 92.15%, Train Loss: 0.2752, Train Accuracy: 90.79%
| Epoch Time: 0m 37s
Epoch [10/20], Validation Loss: 0.2192, Validation Accuracy: 92.73%, Train Loss: 0.2495, Train Accuracy: 91.76%
| Epoch Time: 0m 35s
Epoch [11/20], Validation Loss: 0.2609, Validation Accuracy: 91.67%, Train Loss: 0.2377, Train Accuracy: 92.22%
| Epoch Time: 0m 35s
Epoch [12/20], Validation Loss: 0.1786, Validation Accuracy: 94.57%, Train Loss: 0.2244, Train Accuracy: 92.57%
| Epoch Time: 0m 36s
Epoch [13/20], Validation Loss: 0.2449, Validation Accuracy: 92.25%, Train Loss: 0.1971, Train Accuracy: 93.48%
| Epoch Time: 0m 35s
Epoch [14/20], Validation Loss: 0.1949, Validation Accuracy: 93.76%, Train Loss: 0.1991, Train Accuracy: 93.49%
| Epoch Time: 0m 36s
Epoch [15/20], Validation Loss: 0.1860, Validation Accuracy: 94.15%, Train Loss: 0.1792, Train Accuracy: 93.99%
| Epoch Time: 0m 36s
Epoch [16/20], Validation Loss: 0.2143, Validation Accuracy: 93.04%, Train Loss: 0.1726, Train Accuracy: 94.32%
| Epoch Time: 0m 35s
Epoch [17/20], Validation Loss: 0.1831, Validation Accuracy: 93.80%, Train Loss: 0.1748, Train Accuracy: 94.32%
| Epoch Time: 0m 35s
Epoch [18/20], Validation Loss: 0.1868, Validation Accuracy: 94.46%, Train Loss: 0.1538, Train Accuracy: 94.87%
| Epoch Time: 0m 35s
Epoch [19/20], Validation Loss: 0.1580, Validation Accuracy: 94.89%, Train Loss: 0.1540, Train Accuracy: 94.85%
```

```
| Epoch Time: 0m 36s
Epoch [20/20], Validation Loss: 0.1504, Validation Accuracy: 95.40%, Train Loss:
0.1398, Train Accuracy: 95.41%
| Epoch Time: 0m 36s
Training model for configuration 2...
Epoch [1/20], Validation Loss: 3.3230, Validation Accuracy: 9.47%, Train Loss:
3.6765, Train Accuracy: 9.86%
| Epoch Time: 0m 36s
Epoch [2/20], Validation Loss: 3.3161, Validation Accuracy: 10.71%, Train Loss:
3.3176, Train Accuracy: 9.98%
| Epoch Time: 0m 36s
Epoch [3/20], Validation Loss: 3.3154, Validation Accuracy: 10.71%, Train Loss:
3.3157, Train Accuracy: 10.20%
| Epoch Time: 0m 36s
Epoch [4/20], Validation Loss: 3.3154, Validation Accuracy: 10.71%, Train Loss:
3.3149, Train Accuracy: 10.21%
| Epoch Time: 0m 36s
Epoch [5/20], Validation Loss: 3.3140, Validation Accuracy: 10.71%, Train Loss:
3.3144, Train Accuracy: 9.97%
| Epoch Time: 0m 36s
Epoch [6/20], Validation Loss: 3.3163, Validation Accuracy: 10.03%, Train Loss:
3.3144, Train Accuracy: 9.92%
| Epoch Time: 0m 36s
Epoch [7/20], Validation Loss: 3.3144, Validation Accuracy: 10.71%, Train Loss:
3.3142, Train Accuracy: 10.28%
| Epoch Time: 0m 36s
Epoch [8/20], Validation Loss: 3.3154, Validation Accuracy: 10.03%, Train Loss:
3.3142, Train Accuracy: 10.24%
| Epoch Time: 0m 36s
Epoch [9/20], Validation Loss: 3.3142, Validation Accuracy: 10.03%, Train Loss:
3.3142, Train Accuracy: 10.03%
| Epoch Time: 0m 36s
Epoch [10/20], Validation Loss: 3.3151, Validation Accuracy: 10.71%, Train Loss:
3.3140, Train Accuracy: 10.23%
| Epoch Time: 0m 36s
Epoch [11/20], Validation Loss: 3.3145, Validation Accuracy: 10.71%, Train Loss:
3.3140, Train Accuracy: 10.17%
| Epoch Time: 0m 36s
Epoch [12/20], Validation Loss: 3.3129, Validation Accuracy: 10.71%, Train Loss:
3.3136, Train Accuracy: 10.18%
| Epoch Time: 0m 35s
Epoch [13/20], Validation Loss: 3.3157, Validation Accuracy: 10.71%, Train Loss:
3.3131, Train Accuracy: 10.22%
| Epoch Time: 0m 35s
Epoch [14/20], Validation Loss: 3.3147, Validation Accuracy: 10.03%, Train Loss:
3.3131, Train Accuracy: 10.60%
| Epoch Time: 0m 35s
Epoch [15/20], Validation Loss: 3.3155, Validation Accuracy: 10.03%, Train Loss:
```

```
3.3134, Train Accuracy: 10.30%
| Epoch Time: 0m 36s
Epoch [16/20], Validation Loss: 3.3151, Validation Accuracy: 10.71%, Train Loss:
3.3138, Train Accuracy: 10.26%
| Epoch Time: 0m 36s
Epoch [17/20], Validation Loss: 3.3136, Validation Accuracy: 10.71%, Train Loss:
3.3134, Train Accuracy: 10.24%
| Epoch Time: 0m 36s
Epoch [18/20], Validation Loss: 3.3141, Validation Accuracy: 10.71%, Train Loss:
3.3132, Train Accuracy: 10.24%
| Epoch Time: 0m 36s
Epoch [19/20], Validation Loss: 3.3144, Validation Accuracy: 10.03%, Train Loss:
3.3137, Train Accuracy: 10.04%
| Epoch Time: 0m 36s
Epoch [20/20], Validation Loss: 3.3160, Validation Accuracy: 9.47%, Train Loss:
3.3131, Train Accuracy: 10.33%
| Epoch Time: 0m 36s
Training model for configuration 3...
Epoch [1/20], Validation Loss: 1.9452, Validation Accuracy: 48.54%, Train Loss:
2.5022, Train Accuracy: 34.58%
| Epoch Time: 0m 36s
Epoch [2/20], Validation Loss: 1.4404, Validation Accuracy: 59.39%, Train Loss:
1.7583, Train Accuracy: 51.33%
| Epoch Time: 0m 36s
Epoch [3/20], Validation Loss: 1.2426, Validation Accuracy: 63.55%, Train Loss:
1.4628, Train Accuracy: 58.43%
| Epoch Time: 0m 35s
Epoch [4/20], Validation Loss: 1.0496, Validation Accuracy: 69.05%, Train Loss:
1.2395, Train Accuracy: 64.10%
| Epoch Time: 0m 37s
Epoch [5/20], Validation Loss: 0.8843, Validation Accuracy: 73.70%, Train Loss:
1.0881, Train Accuracy: 67.79%
| Epoch Time: 0m 35s
Epoch [6/20], Validation Loss: 0.7984, Validation Accuracy: 75.98%, Train Loss:
0.9645, Train Accuracy: 71.07%
| Epoch Time: 0m 36s
Epoch [7/20], Validation Loss: 0.7193, Validation Accuracy: 77.89%, Train Loss:
0.8608, Train Accuracy: 74.01%
| Epoch Time: 0m 35s
Epoch [8/20], Validation Loss: 0.6875, Validation Accuracy: 78.98%, Train Loss:
0.7938, Train Accuracy: 75.81%
| Epoch Time: 0m 35s
Epoch [9/20], Validation Loss: 0.6055, Validation Accuracy: 81.10%, Train Loss:
0.7270, Train Accuracy: 77.71%
| Epoch Time: 0m 35s
Epoch [10/20], Validation Loss: 0.5954, Validation Accuracy: 81.49%, Train Loss:
0.6852, Train Accuracy: 78.95%
| Epoch Time: 0m 36s
```

```

Epoch [11/20], Validation Loss: 0.5309, Validation Accuracy: 83.38%, Train Loss: 0.6278, Train Accuracy: 80.34%
| Epoch Time: 0m 36s
Epoch [12/20], Validation Loss: 0.4790, Validation Accuracy: 85.16%, Train Loss: 0.5818, Train Accuracy: 81.58%
| Epoch Time: 0m 36s
Epoch [13/20], Validation Loss: 0.4383, Validation Accuracy: 86.50%, Train Loss: 0.5493, Train Accuracy: 82.76%
| Epoch Time: 0m 36s
Epoch [14/20], Validation Loss: 0.4303, Validation Accuracy: 86.10%, Train Loss: 0.5120, Train Accuracy: 83.81%
| Epoch Time: 0m 36s
Epoch [15/20], Validation Loss: 0.4119, Validation Accuracy: 87.15%, Train Loss: 0.4866, Train Accuracy: 84.59%
| Epoch Time: 0m 36s
Epoch [16/20], Validation Loss: 0.4508, Validation Accuracy: 85.86%, Train Loss: 0.4585, Train Accuracy: 85.39%
| Epoch Time: 0m 36s
Epoch [17/20], Validation Loss: 0.4111, Validation Accuracy: 86.93%, Train Loss: 0.4313, Train Accuracy: 86.30%
| Epoch Time: 0m 36s
Epoch [18/20], Validation Loss: 0.3382, Validation Accuracy: 89.26%, Train Loss: 0.4110, Train Accuracy: 86.78%
| Epoch Time: 0m 35s
Epoch [19/20], Validation Loss: 0.3657, Validation Accuracy: 88.41%, Train Loss: 0.3873, Train Accuracy: 87.51%
| Epoch Time: 0m 35s
Epoch [20/20], Validation Loss: 0.3162, Validation Accuracy: 89.77%, Train Loss: 0.3739, Train Accuracy: 88.05%
| Epoch Time: 0m 35s

```

```

[ ]: # Set Seaborn style
sns.set(style="whitegrid")
fig, axs = plt.subplots(2, 2, figsize=(12, 8))

for i, lr in enumerate(learning_rates):
    axs[0, 0].plot(train_losses_list[i], label=f'Train - lr={lr}')
    axs[0, 1].plot(validation_losses_list[i], label=f'Validation - lr={lr}')
    axs[1, 0].plot(train_accuracy_list[i], label=f'Train - lr={lr}')
    axs[1, 1].plot(validation_accuracy_list[i], label=f'Validation - lr={lr}')

axs[0, 0].set_title('Training Loss for different lr')
axs[0, 0].legend()
axs[0, 0].set_xlabel('Epochs')
axs[0, 0].set_ylabel('Loss (Binary Cross Entropy)')

axs[0, 1].set_title('Validation Loss for different lr')

```

```

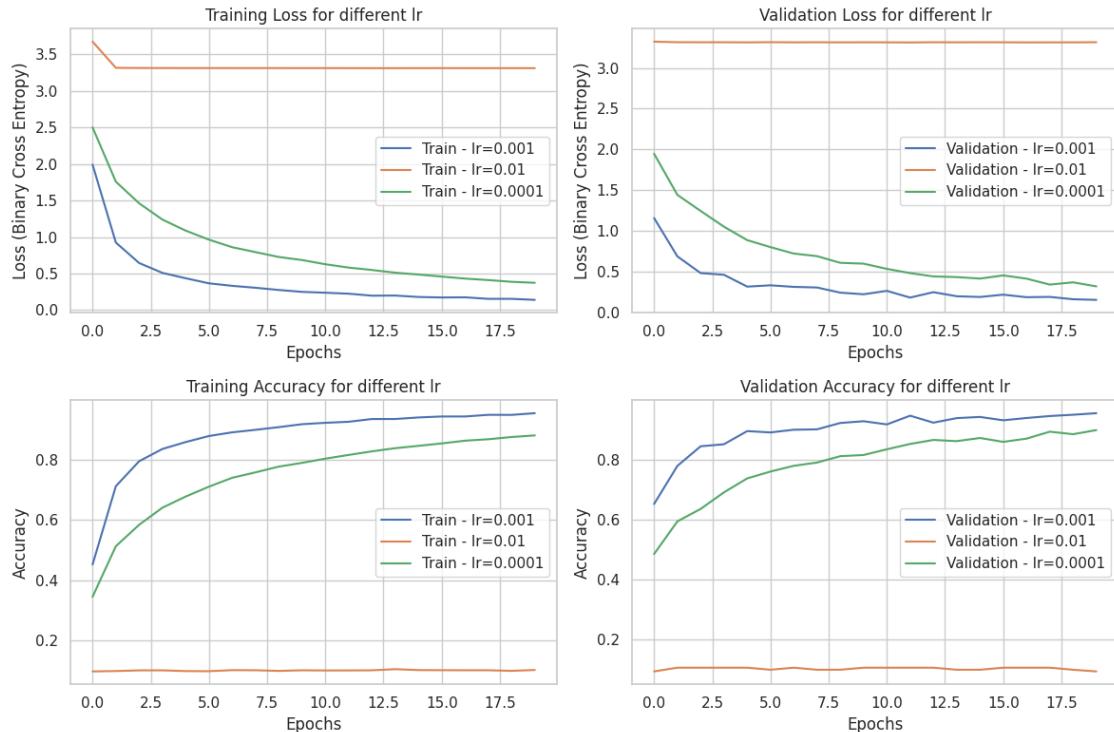
axs[0, 1].legend()
axs[0, 1].set_xlabel('Epochs')
axs[0, 1].set_ylabel('Loss (Binary Cross Entropy)')

axs[1, 0].set_title('Training Accuracy for different lr')
axs[1, 0].legend()
axs[1, 0].set_xlabel('Epochs')
axs[1, 0].set_ylabel('Accuracy')

axs[1, 1].set_title('Validation Accuracy for different lr')
axs[1, 1].legend()
axs[1, 1].set_xlabel('Epochs')
axs[1, 1].set_ylabel('Accuracy')

plt.tight_layout()
plt.show()

```



The best learning among those which are tested is 10^{-3} .

Optimizer selection

```
[ ]: num_epochs = 18
learning_rate = 1e-3
optimizers = ['RMSprop', 'Adam', 'SGD']
```

```

# Lists to store results for each learning rate
train_losses_list = []
validation_losses_list = []
train_accuracy_list = []
validation_accuracy_list = []

for idx, optimizer_name in enumerate(optimizers):
    network = CNN().to(device)
    criterion = nn.CrossEntropyLoss()
    print(f"Training model for configuration {idx+ 1}...")
    if optimizer_name == 'Adam':
        optimizer = optim.Adam(network.parameters(), lr=learning_rate)
    elif optimizer_name == 'SGD':
        optimizer = optim.SGD(network.parameters(), lr=learning_rate)
    elif optimizer_name == 'RMSprop':
        optimizer = optim.RMSprop(network.parameters(), lr=learning_rate)

    # Train the model
    train_avg_loss, validation_avg_loss, validation_accuracy, train_accuracy =
    ↪train(network, num_epochs, trainloader, validationloader, criterion,
    ↪optimizer, validation_phase=True)

    # Store results for plotting
    train_losses_list.append(train_avg_loss)
    validation_losses_list.append(validation_avg_loss)
    train_accuracy_list.append(train_accuracy)
    validation_accuracy_list.append(validation_accuracy)

```

Training model for configuration 1...

Epoch [1/18], Validation Loss: 2.3102, Validation Accuracy: 37.93%, Train Loss: 2.6718, Train Accuracy: 29.80%

| Epoch Time: 0m 35s

Epoch [2/18], Validation Loss: 1.6797, Validation Accuracy: 54.92%, Train Loss: 1.4774, Train Accuracy: 56.45%

| Epoch Time: 0m 35s

Epoch [3/18], Validation Loss: 1.5975, Validation Accuracy: 61.59%, Train Loss: 0.9859, Train Accuracy: 69.38%

| Epoch Time: 0m 35s

Epoch [4/18], Validation Loss: 0.9087, Validation Accuracy: 73.72%, Train Loss: 0.7607, Train Accuracy: 75.72%

| Epoch Time: 0m 34s

Epoch [5/18], Validation Loss: 0.9755, Validation Accuracy: 73.48%, Train Loss: 0.6256, Train Accuracy: 79.81%

| Epoch Time: 0m 35s

Epoch [6/18], Validation Loss: 1.4048, Validation Accuracy: 65.20%, Train Loss: 0.5347, Train Accuracy: 82.70%

| Epoch Time: 0m 34s

```
Epoch [7/18], Validation Loss: 0.6019, Validation Accuracy: 81.35%, Train Loss: 0.4708, Train Accuracy: 84.70%
| Epoch Time: 0m 37s
Epoch [8/18], Validation Loss: 0.6055, Validation Accuracy: 81.69%, Train Loss: 0.4285, Train Accuracy: 86.01%
| Epoch Time: 0m 35s
Epoch [9/18], Validation Loss: 0.4942, Validation Accuracy: 85.42%, Train Loss: 0.3859, Train Accuracy: 87.28%
| Epoch Time: 0m 34s
Epoch [10/18], Validation Loss: 0.5562, Validation Accuracy: 83.24%, Train Loss: 0.3496, Train Accuracy: 88.44%
| Epoch Time: 0m 35s
Epoch [11/18], Validation Loss: 1.0359, Validation Accuracy: 73.69%, Train Loss: 0.3274, Train Accuracy: 89.16%
| Epoch Time: 0m 34s
Epoch [12/18], Validation Loss: 0.4902, Validation Accuracy: 85.51%, Train Loss: 0.3056, Train Accuracy: 89.94%
| Epoch Time: 0m 35s
Epoch [13/18], Validation Loss: 0.5294, Validation Accuracy: 85.52%, Train Loss: 0.2830, Train Accuracy: 90.64%
| Epoch Time: 0m 35s
Epoch [14/18], Validation Loss: 0.5655, Validation Accuracy: 83.84%, Train Loss: 0.2696, Train Accuracy: 91.16%
| Epoch Time: 0m 36s
Epoch [15/18], Validation Loss: 0.5261, Validation Accuracy: 85.15%, Train Loss: 0.2520, Train Accuracy: 91.72%
| Epoch Time: 0m 34s
Epoch [16/18], Validation Loss: 0.5772, Validation Accuracy: 83.06%, Train Loss: 0.2376, Train Accuracy: 92.26%
| Epoch Time: 0m 37s
Epoch [17/18], Validation Loss: 0.7734, Validation Accuracy: 81.81%, Train Loss: 0.2309, Train Accuracy: 92.44%
| Epoch Time: 0m 36s
Epoch [18/18], Validation Loss: 0.5326, Validation Accuracy: 86.33%, Train Loss: 0.2181, Train Accuracy: 92.92%
| Epoch Time: 0m 36s
Training model for configuration 2...
Epoch [1/18], Validation Loss: 1.9999, Validation Accuracy: 52.63%, Train Loss: 2.1795, Train Accuracy: 40.28%
| Epoch Time: 0m 39s
Epoch [2/18], Validation Loss: 1.1150, Validation Accuracy: 66.76%, Train Loss: 1.1533, Train Accuracy: 65.06%
| Epoch Time: 0m 35s
Epoch [3/18], Validation Loss: 1.0292, Validation Accuracy: 71.87%, Train Loss: 0.7842, Train Accuracy: 75.41%
| Epoch Time: 0m 36s
Epoch [4/18], Validation Loss: 0.8582, Validation Accuracy: 76.11%, Train Loss: 0.6181, Train Accuracy: 80.14%
```

```
| Epoch Time: 0m 36s
Epoch [5/18], Validation Loss: 0.7083, Validation Accuracy: 79.90%, Train Loss: 0.5239, Train Accuracy: 82.91%
| Epoch Time: 0m 35s
Epoch [6/18], Validation Loss: 0.6718, Validation Accuracy: 80.20%, Train Loss: 0.4561, Train Accuracy: 85.04%
| Epoch Time: 0m 36s
Epoch [7/18], Validation Loss: 0.5657, Validation Accuracy: 83.56%, Train Loss: 0.4130, Train Accuracy: 86.45%
| Epoch Time: 0m 36s
Epoch [8/18], Validation Loss: 0.5004, Validation Accuracy: 85.37%, Train Loss: 0.3708, Train Accuracy: 87.89%
| Epoch Time: 0m 37s
Epoch [9/18], Validation Loss: 0.5473, Validation Accuracy: 85.23%, Train Loss: 0.3317, Train Accuracy: 89.12%
| Epoch Time: 0m 34s
Epoch [10/18], Validation Loss: 0.5339, Validation Accuracy: 84.95%, Train Loss: 0.3098, Train Accuracy: 89.77%
| Epoch Time: 0m 36s
Epoch [11/18], Validation Loss: 0.5063, Validation Accuracy: 84.94%, Train Loss: 0.2920, Train Accuracy: 90.33%
| Epoch Time: 0m 36s
Epoch [12/18], Validation Loss: 0.5192, Validation Accuracy: 84.63%, Train Loss: 0.2716, Train Accuracy: 91.02%
| Epoch Time: 0m 35s
Epoch [13/18], Validation Loss: 0.4544, Validation Accuracy: 87.11%, Train Loss: 0.2581, Train Accuracy: 91.59%
| Epoch Time: 0m 36s
Epoch [14/18], Validation Loss: 0.4790, Validation Accuracy: 87.22%, Train Loss: 0.2458, Train Accuracy: 91.87%
| Epoch Time: 0m 37s
Epoch [15/18], Validation Loss: 0.4393, Validation Accuracy: 87.36%, Train Loss: 0.2346, Train Accuracy: 92.35%
| Epoch Time: 0m 35s
Epoch [16/18], Validation Loss: 0.4313, Validation Accuracy: 87.97%, Train Loss: 0.2154, Train Accuracy: 92.90%
| Epoch Time: 0m 39s
Epoch [17/18], Validation Loss: 0.4215, Validation Accuracy: 88.12%, Train Loss: 0.2161, Train Accuracy: 92.83%
| Epoch Time: 0m 38s
Epoch [18/18], Validation Loss: 0.4281, Validation Accuracy: 87.98%, Train Loss: 0.1973, Train Accuracy: 93.45%
| Epoch Time: 0m 38s
Training model for configuration 3...
Epoch [1/18], Validation Loss: 3.6395, Validation Accuracy: 2.55%, Train Loss: 3.6493, Train Accuracy: 2.68%
| Epoch Time: 0m 36s
Epoch [2/18], Validation Loss: 3.6201, Validation Accuracy: 2.55%, Train Loss:
```

```
3.6305, Train Accuracy: 2.68%
| Epoch Time: 0m 36s
Epoch [3/18], Validation Loss: 3.5921, Validation Accuracy: 6.89%, Train Loss:
3.6074, Train Accuracy: 3.23%
| Epoch Time: 0m 37s
Epoch [4/18], Validation Loss: 3.5153, Validation Accuracy: 9.52%, Train Loss:
3.5595, Train Accuracy: 8.92%
| Epoch Time: 0m 35s
Epoch [5/18], Validation Loss: 3.3889, Validation Accuracy: 17.22%, Train Loss:
3.4433, Train Accuracy: 9.98%
| Epoch Time: 0m 35s
Epoch [6/18], Validation Loss: 3.3182, Validation Accuracy: 22.18%, Train Loss:
3.3477, Train Accuracy: 18.17%
| Epoch Time: 0m 36s
Epoch [7/18], Validation Loss: 3.2527, Validation Accuracy: 23.69%, Train Loss:
3.2914, Train Accuracy: 20.28%
| Epoch Time: 0m 35s
Epoch [8/18], Validation Loss: 3.1621, Validation Accuracy: 23.78%, Train Loss:
3.2311, Train Accuracy: 20.19%
| Epoch Time: 0m 37s
Epoch [9/18], Validation Loss: 3.0303, Validation Accuracy: 25.68%, Train Loss:
3.1442, Train Accuracy: 20.84%
| Epoch Time: 0m 37s
Epoch [10/18], Validation Loss: 2.9054, Validation Accuracy: 26.69%, Train Loss:
3.0319, Train Accuracy: 23.47%
| Epoch Time: 0m 36s
Epoch [11/18], Validation Loss: 2.8026, Validation Accuracy: 28.96%, Train Loss:
2.9255, Train Accuracy: 25.03%
| Epoch Time: 0m 36s
Epoch [12/18], Validation Loss: 2.7247, Validation Accuracy: 29.75%, Train Loss:
2.8275, Train Accuracy: 27.10%
| Epoch Time: 0m 36s
Epoch [13/18], Validation Loss: 2.6839, Validation Accuracy: 29.28%, Train Loss:
2.7424, Train Accuracy: 28.58%
| Epoch Time: 0m 36s
Epoch [14/18], Validation Loss: 2.5792, Validation Accuracy: 31.87%, Train Loss:
2.6698, Train Accuracy: 30.02%
| Epoch Time: 0m 35s
Epoch [15/18], Validation Loss: 2.5584, Validation Accuracy: 30.98%, Train Loss:
2.5974, Train Accuracy: 31.36%
| Epoch Time: 0m 36s
Epoch [16/18], Validation Loss: 2.5178, Validation Accuracy: 32.27%, Train Loss:
2.5271, Train Accuracy: 32.39%
| Epoch Time: 0m 36s
Epoch [17/18], Validation Loss: 2.3955, Validation Accuracy: 35.14%, Train Loss:
2.4717, Train Accuracy: 33.48%
| Epoch Time: 0m 35s
Epoch [18/18], Validation Loss: 2.3846, Validation Accuracy: 36.62%, Train Loss:
```

```
2.4202, Train Accuracy: 34.53%
| Epoch Time: 0m 35s
```

```
[ ]: sns.set(style="whitegrid")

fig, axs = plt.subplots(2, 2, figsize=(12, 8))

for i, optimizer_name in enumerate(optimizers):
    axs[0, 0].plot(train_losses_list[i], label=f'Train - {optimizer_name}')
    axs[0, 1].plot(validation_losses_list[i], label=f'Validation - {optimizer_name}')
    axs[1, 0].plot(train_accuracy_list[i], label=f'Train - {optimizer_name}')
    axs[1, 1].plot(validation_accuracy_list[i], label=f'Validation - {optimizer_name}')

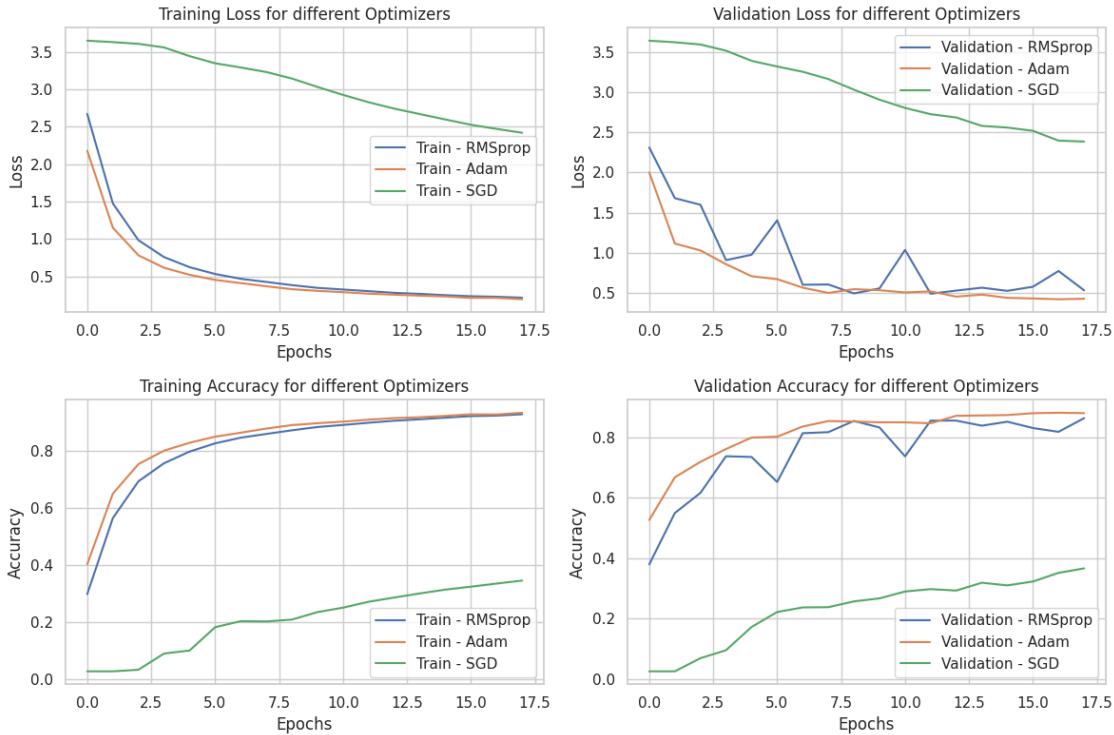
axs[0, 0].set_title('Training Loss for different Optimizers')
axs[0, 0].legend()
axs[0, 0].set_xlabel('Epochs')
axs[0, 0].set_ylabel('Loss')

axs[0, 1].set_title('Validation Loss for different Optimizers')
axs[0, 1].legend()
axs[0, 1].set_xlabel('Epochs')
axs[0, 1].set_ylabel('Loss')

axs[1, 0].set_title('Training Accuracy for different Optimizers')
axs[1, 0].legend()
axs[1, 0].set_xlabel('Epochs')
axs[1, 0].set_ylabel('Accuracy')

axs[1, 1].set_title('Validation Accuracy for different Optimizers')
axs[1, 1].legend()
axs[1, 1].set_xlabel('Epochs')
axs[1, 1].set_ylabel('Accuracy')

plt.tight_layout()
plt.show()
```



3 D. 1) Train the model over the color images

```
[ ]: # Define hyperparameters
learning_rate = 1e-3
num_epochs = 25

network = CNN(grayscale=False).to(device)
network.apply(initialize_parameters)
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(network.parameters(), lr=learning_rate)

print(f'The model has {count_parameters(network)} trainable parameters')
```

The model has 6,846,758 trainable parameters

```
[ ]: # Train the model
train_avg_loss, validation_avg_loss, validation_accuracy, train_accuracy =
    ↪train(network, num_epochs, trainloader, validationloader, criterion,
    ↪optimizer, validation_phase=True)
```

Epoch [1/25], Validation Loss: 1.2302, Validation Accuracy: 64.63%, Train Loss: 2.3124, Train Accuracy: 44.40%
| Epoch Time: 0m 36s

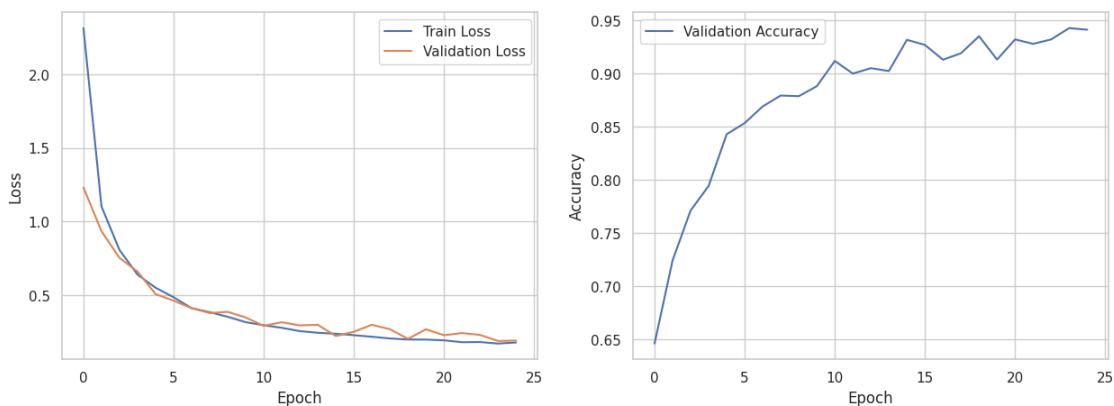
```
Epoch [2/25], Validation Loss: 0.9330, Validation Accuracy: 72.44%, Train Loss: 1.1015, Train Accuracy: 67.76%
| Epoch Time: 0m 37s
Epoch [3/25], Validation Loss: 0.7526, Validation Accuracy: 77.12%, Train Loss: 0.8058, Train Accuracy: 75.54%
| Epoch Time: 0m 37s
Epoch [4/25], Validation Loss: 0.6592, Validation Accuracy: 79.43%, Train Loss: 0.6387, Train Accuracy: 80.11%
| Epoch Time: 0m 36s
Epoch [5/25], Validation Loss: 0.5066, Validation Accuracy: 84.28%, Train Loss: 0.5500, Train Accuracy: 82.81%
| Epoch Time: 0m 36s
Epoch [6/25], Validation Loss: 0.4616, Validation Accuracy: 85.33%, Train Loss: 0.4851, Train Accuracy: 84.62%
| Epoch Time: 0m 35s
Epoch [7/25], Validation Loss: 0.4115, Validation Accuracy: 86.89%, Train Loss: 0.4111, Train Accuracy: 86.85%
| Epoch Time: 0m 35s
Epoch [8/25], Validation Loss: 0.3779, Validation Accuracy: 87.92%, Train Loss: 0.3836, Train Accuracy: 87.96%
| Epoch Time: 0m 36s
Epoch [9/25], Validation Loss: 0.3860, Validation Accuracy: 87.85%, Train Loss: 0.3526, Train Accuracy: 88.75%
| Epoch Time: 0m 36s
Epoch [10/25], Validation Loss: 0.3478, Validation Accuracy: 88.79%, Train Loss: 0.3158, Train Accuracy: 90.00%
| Epoch Time: 0m 37s
Epoch [11/25], Validation Loss: 0.2908, Validation Accuracy: 91.16%, Train Loss: 0.2952, Train Accuracy: 90.59%
| Epoch Time: 0m 36s
Epoch [12/25], Validation Loss: 0.3157, Validation Accuracy: 89.97%, Train Loss: 0.2776, Train Accuracy: 91.14%
| Epoch Time: 0m 37s
Epoch [13/25], Validation Loss: 0.2940, Validation Accuracy: 90.48%, Train Loss: 0.2554, Train Accuracy: 91.79%
| Epoch Time: 0m 37s
Epoch [14/25], Validation Loss: 0.2978, Validation Accuracy: 90.21%, Train Loss: 0.2438, Train Accuracy: 92.27%
| Epoch Time: 0m 36s
Epoch [15/25], Validation Loss: 0.2218, Validation Accuracy: 93.14%, Train Loss: 0.2371, Train Accuracy: 92.28%
| Epoch Time: 0m 37s
Epoch [16/25], Validation Loss: 0.2509, Validation Accuracy: 92.66%, Train Loss: 0.2276, Train Accuracy: 92.88%
| Epoch Time: 0m 39s
Epoch [17/25], Validation Loss: 0.2982, Validation Accuracy: 91.28%, Train Loss: 0.2169, Train Accuracy: 93.09%
| Epoch Time: 0m 37s
```

```

Epoch [18/25], Validation Loss: 0.2682, Validation Accuracy: 91.88%, Train Loss: 0.2055, Train Accuracy: 93.53%
| Epoch Time: 0m 38s
Epoch [19/25], Validation Loss: 0.2024, Validation Accuracy: 93.48%, Train Loss: 0.1981, Train Accuracy: 93.70%
| Epoch Time: 0m 38s
Epoch [20/25], Validation Loss: 0.2676, Validation Accuracy: 91.29%, Train Loss: 0.1977, Train Accuracy: 93.80%
| Epoch Time: 0m 37s
Epoch [21/25], Validation Loss: 0.2273, Validation Accuracy: 93.18%, Train Loss: 0.1927, Train Accuracy: 94.00%
| Epoch Time: 0m 38s
Epoch [22/25], Validation Loss: 0.2419, Validation Accuracy: 92.77%, Train Loss: 0.1798, Train Accuracy: 94.32%
| Epoch Time: 0m 37s
Epoch [23/25], Validation Loss: 0.2293, Validation Accuracy: 93.18%, Train Loss: 0.1811, Train Accuracy: 94.18%
| Epoch Time: 0m 38s
Epoch [24/25], Validation Loss: 0.1881, Validation Accuracy: 94.24%, Train Loss: 0.1706, Train Accuracy: 94.63%
| Epoch Time: 0m 38s
Epoch [25/25], Validation Loss: 0.1911, Validation Accuracy: 94.09%, Train Loss: 0.1780, Train Accuracy: 94.52%
| Epoch Time: 0m 37s

```

```
[ ]: plot_training_results(train_avg_loss, validation_avg_loss, validation_accuracy, is_validation=True)
```



```
[ ]: # Combining train and validation datasets for testing the model
combined_train_set = ConcatDataset([train_set, validation_set])
combined_train_loader = torch.utils.data.DataLoader(combined_train_set, batch_size=64, shuffle=True, num_workers=2)
```

```
[ ]: # Test the model
network = CNN(grayscale=False).to(device)
network.apply(initialize_parameters)
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(network.parameters(), lr=learning_rate)

train_avg_loss, test_avg_loss, test_accuracy, train_accuracy = train(network, □
    ↳num_epochs, combined_train_loader, testloader, criterion, □
    ↳optimizer, validation_phase=False)
```

Epoch [1/25], Test Loss: 1.6582, Test Accuracy: 54.94%, Train Loss: 2.1497,
 Train Accuracy: 46.78%
 | Epoch Time: 0m 50s
 Epoch [2/25], Test Loss: 1.0267, Test Accuracy: 70.14%, Train Loss: 1.0515,
 Train Accuracy: 68.69%
 | Epoch Time: 0m 49s
 Epoch [3/25], Test Loss: 0.8566, Test Accuracy: 73.95%, Train Loss: 0.7440,
 Train Accuracy: 76.77%
 | Epoch Time: 0m 48s
 Epoch [4/25], Test Loss: 0.8391, Test Accuracy: 75.17%, Train Loss: 0.5684,
 Train Accuracy: 82.01%
 | Epoch Time: 0m 47s
 Epoch [5/25], Test Loss: 0.7107, Test Accuracy: 78.57%, Train Loss: 0.5218,
 Train Accuracy: 83.45%
 | Epoch Time: 0m 49s
 Epoch [6/25], Test Loss: 0.6762, Test Accuracy: 79.89%, Train Loss: 0.4537,
 Train Accuracy: 85.73%
 | Epoch Time: 0m 47s
 Epoch [7/25], Test Loss: 0.5481, Test Accuracy: 84.00%, Train Loss: 0.3987,
 Train Accuracy: 86.85%
 | Epoch Time: 0m 49s
 Epoch [8/25], Test Loss: 0.5222, Test Accuracy: 84.63%, Train Loss: 0.3486,
 Train Accuracy: 88.63%
 | Epoch Time: 0m 49s
 Epoch [9/25], Test Loss: 0.4841, Test Accuracy: 85.80%, Train Loss: 0.3208,
 Train Accuracy: 89.67%
 | Epoch Time: 0m 48s
 Epoch [10/25], Test Loss: 0.6469, Test Accuracy: 83.64%, Train Loss: 0.3027,
 Train Accuracy: 90.18%
 | Epoch Time: 0m 49s
 Epoch [11/25], Test Loss: 0.5310, Test Accuracy: 84.90%, Train Loss: 0.2727,
 Train Accuracy: 91.07%
 | Epoch Time: 0m 48s
 Epoch [12/25], Test Loss: 0.5133, Test Accuracy: 86.38%, Train Loss: 0.2838,
 Train Accuracy: 90.83%
 | Epoch Time: 0m 48s
 Epoch [13/25], Test Loss: 0.4769, Test Accuracy: 86.87%, Train Loss: 0.2360,
 Train Accuracy: 92.14%

```
| Epoch Time: 0m 48s
Epoch [14/25], Test Loss: 0.5146, Test Accuracy: 85.94%, Train Loss: 0.2309,
Train Accuracy: 92.55%
| Epoch Time: 0m 48s
Epoch [15/25], Test Loss: 0.4240, Test Accuracy: 88.20%, Train Loss: 0.2224,
Train Accuracy: 92.79%
| Epoch Time: 0m 49s
Epoch [16/25], Test Loss: 0.4716, Test Accuracy: 86.90%, Train Loss: 0.2145,
Train Accuracy: 93.08%
| Epoch Time: 0m 48s
Epoch [17/25], Test Loss: 0.5401, Test Accuracy: 86.50%, Train Loss: 0.2297,
Train Accuracy: 92.73%
| Epoch Time: 0m 49s
Epoch [18/25], Test Loss: 0.4163, Test Accuracy: 89.00%, Train Loss: 0.1935,
Train Accuracy: 93.84%
| Epoch Time: 0m 48s
Epoch [19/25], Test Loss: 0.4402, Test Accuracy: 88.19%, Train Loss: 0.1776,
Train Accuracy: 94.36%
| Epoch Time: 0m 48s
Epoch [20/25], Test Loss: 0.4542, Test Accuracy: 88.63%, Train Loss: 0.1766,
Train Accuracy: 94.35%
| Epoch Time: 0m 49s
Epoch [21/25], Test Loss: 0.4572, Test Accuracy: 88.48%, Train Loss: 0.1696,
Train Accuracy: 94.56%
| Epoch Time: 0m 49s
Epoch [22/25], Test Loss: 0.3810, Test Accuracy: 89.63%, Train Loss: 0.1622,
Train Accuracy: 94.92%
| Epoch Time: 0m 49s
Epoch [23/25], Test Loss: 0.3951, Test Accuracy: 88.98%, Train Loss: 0.1627,
Train Accuracy: 94.83%
| Epoch Time: 0m 48s
Epoch [24/25], Test Loss: 0.6739, Test Accuracy: 85.69%, Train Loss: 0.1595,
Train Accuracy: 95.06%
| Epoch Time: 0m 48s
Epoch [25/25], Test Loss: 0.3686, Test Accuracy: 90.00%, Train Loss: 0.1507,
Train Accuracy: 95.26%
| Epoch Time: 0m 49s
```

```
[ ]: plot_training_results(train_avg_loss, test_avg_loss, test_accuracy,  
                           ↵is_validation=False)
```



Examining the model

```
[ ]: network.load_state_dict(torch.load('model.pt'))
```

```
[ ]: <All keys matched successfully>
```

```
[26]: def get_predictions(model, iterator, device):
    model.eval()
    model.to(device)

    images = []
    labels = []
    probs = []
    corrects = []

    with torch.no_grad():
        for (x, y) in iterator:
            x = x.to(device)
            y = y.to(device)

            outputs = model(x)
            if isinstance(outputs, tuple):
                y_pred = outputs[0]
            else:
                y_pred = outputs

            y_prob = F.softmax(y_pred, dim=-1)
            _, predicted = torch.max(y_prob, 1)

            images.append(x.cpu())
            labels.append(y.cpu())
            probs.append(y_prob.cpu())
```

```

# Determine correctness of predictions
correct = predicted.eq(y)
corrects.extend(correct.cpu().tolist())

images = torch.cat(images, dim=0)
labels = torch.cat(labels, dim=0)
probs = torch.cat(probs, dim=0)

return images, labels, probs, corrects

```

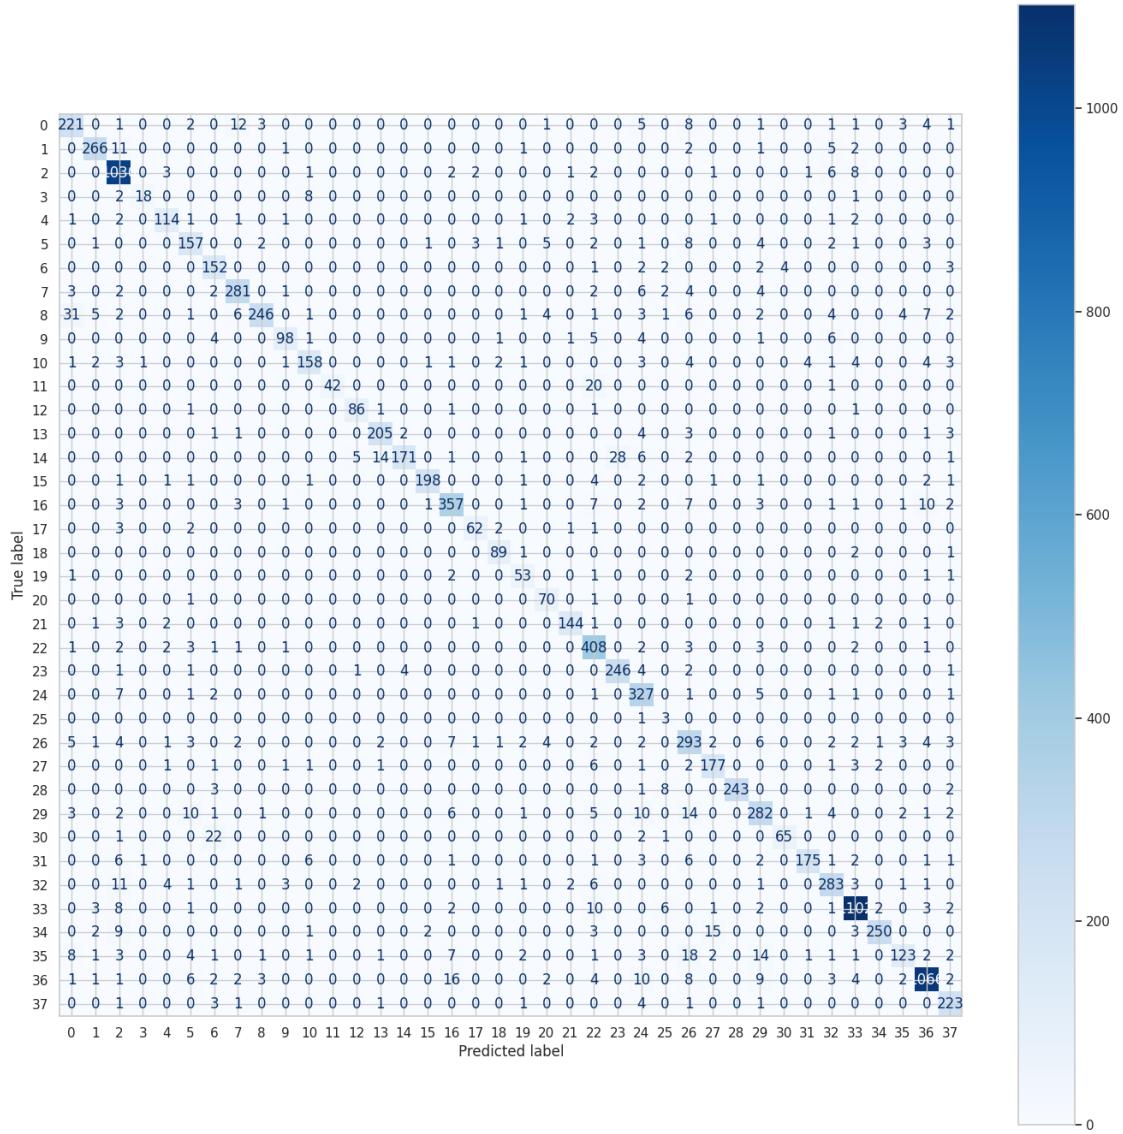
```
[ ]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
images, labels, probs, corrects = get_predictions(network, testloader, device)
pred_labels = torch.argmax(probs, 1)

print(f"There are {len(corrects)} correct predictions.")
```

There are 10562 correct predictions.

```
[27]: # Define the plot_confusion_matrix function
def plot_confusion_matrix(labels, pred_labels):
    fig = plt.figure(figsize=(17, 17))
    ax = fig.add_subplot(1, 1, 1)
    cm = confusion_matrix(labels, pred_labels)
    cm = ConfusionMatrixDisplay(cm, display_labels=range(38))
    cm.plot(values_format='d', cmap='Blues', ax=ax)
```

```
[ ]: plot_confusion_matrix(labels, pred_labels)
```



```
[ ]: # Now, extract and sort incorrect examples
incorrect_examples = []
for image, label, prob, correct in zip(images, labels, probs, corrects):
    if not correct:
        incorrect_examples.append((image, label, prob))

incorrect_examples.sort(reverse=True, key=lambda x: torch.max(x[2], dim=0).
                           values)
print(f"There are {len(incorrect_examples)} incorrect predictions.")
```

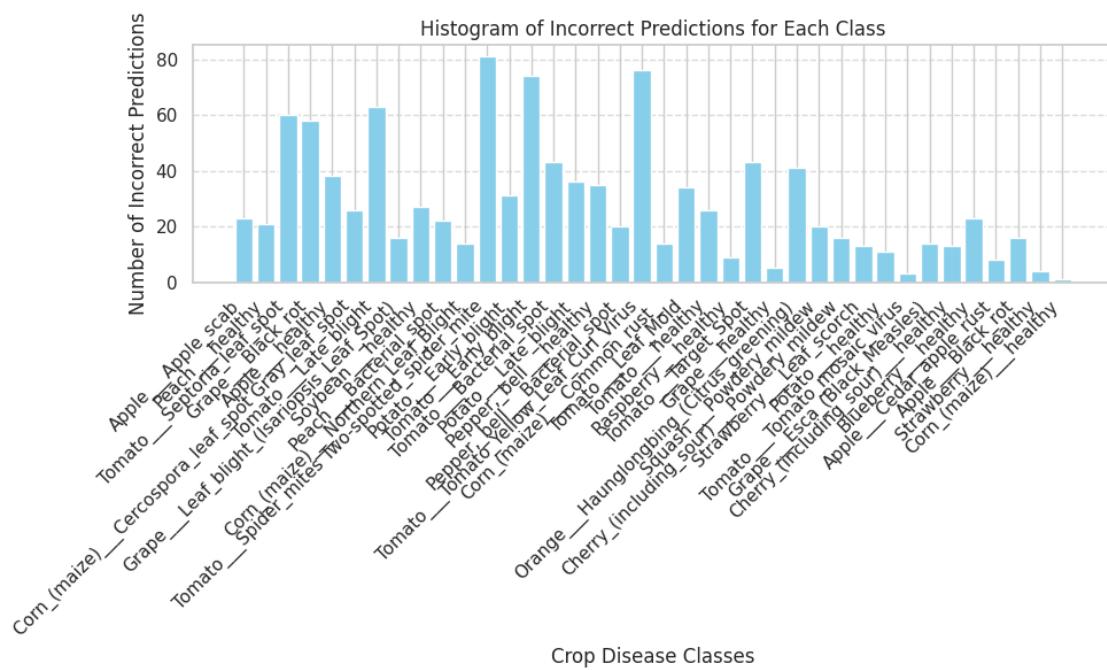
There are 1078 incorrect predictions.

```
[ ]: incorrect_counts = {}

# Count the occurrences of each class label in the incorrect examples
for _, label, _ in incorrect_examples:
    if label.item() not in incorrect_counts:
        incorrect_counts[label.item()] = 1
    else:
        incorrect_counts[label.item()] += 1

class_names = [train_set.classes[label] for label in incorrect_counts.keys()]

# Plotting the histogram
plt.figure(figsize=(10, 6))
plt.bar(class_names, incorrect_counts.values(), color='skyblue')
plt.xlabel('Crop Disease Classes')
plt.ylabel('Number of Incorrect Predictions')
plt.title('Histogram of Incorrect Predictions for Each Class')
plt.xticks(rotation=45, ha='right')
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()
```



```
[ ]: data = {'Class Name': class_names, 'Incorrect Counts': list(incorrect_counts.  
    ↪values())}  
df1 = pd.DataFrame(data)
```

```

df1.head(5)

[ ]:          Class Name  Incorrect Counts
0      Apple___Apple_scab           23
1      Peach___healthy            21
2 Tomato___Septoria_leaf_spot        60
3      Grape___Black_rot           58
4      Apple___healthy             38

[ ]: class_counts = defaultdict(int)
for images, labels in testloader:
    for label in labels:
        class_counts[label.item()] += 1

class_labels = sorted(class_counts.keys())
counts = [class_counts[label] for label in class_labels]
colors = sns.color_palette("husl", len(class_labels))

class_names = [train_set.classes[label] for label in class_labels]

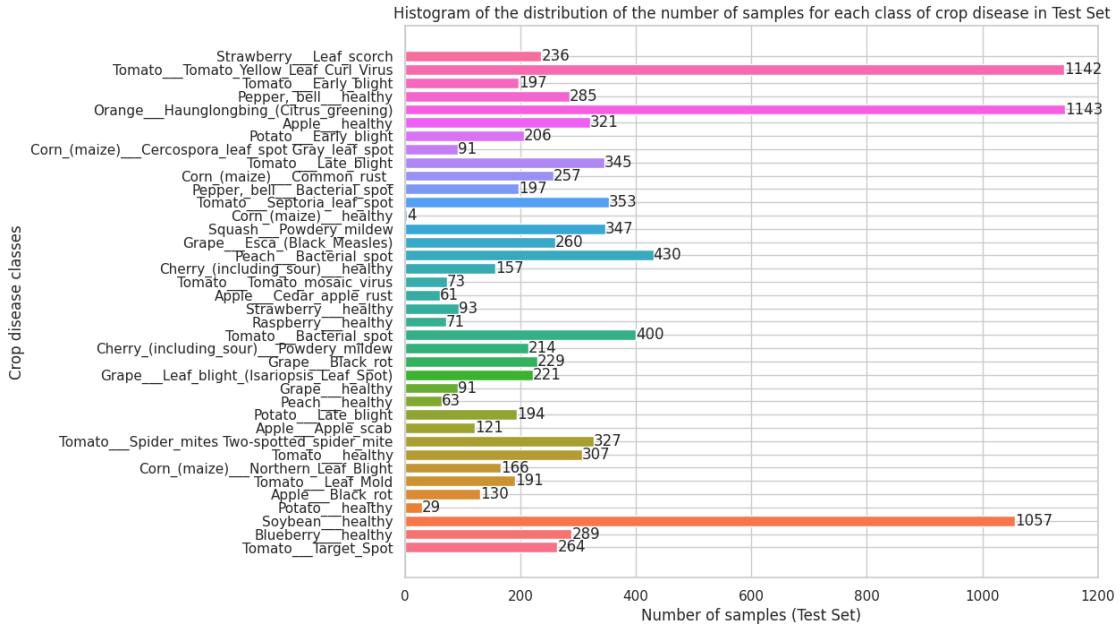
sns.set(style="whitegrid")
plt.figure(figsize=(10, 8))
plt.barh(class_names, counts, color=colors)
plt.xlabel('Number of samples (Test Set)')
plt.ylabel('Crop disease classes')
plt.title('Histogram of the distribution of the number of samples for each class  
of crop disease in Test Set')

for i, (count, name) in enumerate(zip(counts, class_names)):
    plt.text(count, i, str(count), ha='left', va='center')

plt.show()

/usr/lib/python3.10/multiprocessing/popen_fork.py:66: RuntimeWarning: os.fork()  
was called. os.fork() is incompatible with multithreaded code, and JAX is  
multithreaded, so this will likely lead to a deadlock.  
    self.pid = os.fork()  
/usr/lib/python3.10/multiprocessing/popen_fork.py:66: RuntimeWarning: os.fork()  
was called. os.fork() is incompatible with multithreaded code, and JAX is  
multithreaded, so this will likely lead to a deadlock.  
    self.pid = os.fork()

```



```
[ ]: data = {'Class Name': class_names, 'Counts': counts}
df2 = pd.DataFrame(data)
df2.head(5)
```

```
[ ]:      Class Name  Counts
0 Tomato__Target_Spot    264
1 Blueberry__healthy     289
2 Soybean__healthy       1057
3 Potato__healthy        29
4 Apple__Black_rot       130
```

```
[ ]: merged_df = pd.merge(df1, df2, on='Class Name', suffixes=('_incorrect', '_test'))
merged_df.head(3)
```

```
[ ]:      Class Name  Incorrect Counts  Counts
0      Apple__Apple_scab            23    121
1      Peach__healthy              21     63
2 Tomato__Septoria_leaf_spot      60    353
```

```
[ ]: # Calculate the rate of success for each class in percentage
merged_df['Success Rate (%)'] = ((merged_df['Counts'] - merged_df['IncorrectCounts']) / merged_df['Counts']) * 100
merged_df
```

```
[ ]:      Class Name  Incorrect Counts \
0      Apple__Apple_scab            23
```

1	Peach___healthy	21
2	Tomato___Septoria_leaf_spot	60
3	Grape___Black_rot	58
4	Apple___healthy	38
5	Corn_(maize)___Cercospora_leaf_spot_Gray_leaf_...	26
6	Tomato___Late_blight	63
7	Grape___Leaf_blight_(Isariopsis_Leaf_Spot)	16
8	Soybean___healthy	27
9	Peach___Bacterial_spot	22
10	Corn_(maize)___Northern_Leaf_Blight	14
11	Tomato___Spider_mites Two-spotted_spider_mite	81
12	Potato___Early_blight	31
13	Tomato___Early_blight	74
14	Tomato___Bacterial_spot	43
15	Potato___Late_blight	36
16	Pepper,_bell___healthy	35
17	Pepper,_bell___Bacterial_spot	20
18	Tomato___Tomato_Yellow_Leaf_Curl_Virus	76
19	Corn_(maize)___Common_rust_	14
20	Tomato___Leaf_Mold	34
21	Tomato___healthy	26
22	Raspberry___healthy	9
23	Tomato___Target_Spot	43
24	Grape___healthy	5
25	Orange___Haunglongbing_(Citrus_greening)	41
26	Squash___Powdery_mildew	20
27	Cherry_(including_sour)___Powdery_mildew	16
28	Strawberry___Leaf_scorch	13
29	Potato___healthy	11
30	Tomato___Tomato_mosaic_virus	3
31	Grape___Esca_(Black_Measles)	14
32	Cherry_(including_sour)___healthy	13
33	Blueberry___healthy	23
34	Apple___Cedar_apple_rust	8
35	Apple___Black_rot	16
36	Strawberry___healthy	4
37	Corn_(maize)___healthy	1

	Counts	Success Rate (%)
0	121	80.991736
1	63	66.666667
2	353	83.002833
3	229	74.672489
4	321	88.161994
5	91	71.428571
6	345	81.739130
7	221	92.760181

8	1057	97.445601
9	430	94.883721
10	166	91.566265
11	327	75.229358
12	206	84.951456
13	197	62.436548
14	400	89.250000
15	194	81.443299
16	285	87.719298
17	197	89.847716
18	1142	93.345009
19	257	94.552529
20	191	82.198953
21	307	91.530945
22	71	87.323944
23	264	83.712121
24	91	94.505495
25	1143	96.412948
26	347	94.236311
27	214	92.523364
28	236	94.491525
29	29	62.068966
30	73	95.890411
31	260	94.615385
32	157	91.719745
33	289	92.041522
34	61	86.885246
35	130	87.692308
36	93	95.698925
37	4	75.000000

```
[38]: def plot_most_incorrect(incorrect, n_images):
    rows = int(np.sqrt(n_images))
    cols = int(np.sqrt(n_images))

    fig = plt.figure(figsize=(20, 10))
    for i in range(rows*cols):
        ax = fig.add_subplot(rows, cols, i+1)
        image, true_label, probs = incorrect[i]
        true_prob = probs[true_label]
        incorrect_prob, incorrect_label = torch.max(probs, dim=0)

        # Normalize image tensor to [0, 1]
        image = image.permute(1, 2, 0).cpu().numpy()
        image_min, image_max = np.min(image), np.max(image)
        image = (image - image_min) / (image_max - image_min)
```

```

    ax.imshow(image, cmap='bone')
    ax.set_title(f'true label: {true_label} ({true_prob:.3f})\n'
                 f'pred label: {incorrect_label.item()} ({incorrect_prob:.3f})')
    ax.axis('off')
fig.subplots_adjust(hspace=0.5)

```

[]: N_IMAGES = 25

plot_most_incorrect(incorrect_examples, N_IMAGES)



[]: !pip install tqdm

Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages (4.66.2)

[24]: from tqdm import tqdm

```

def get_representations(model, iterator, device):

    model.eval()

    outputs = []
    intermediates = []
    labels = []

    with torch.no_grad():

```

```

for (x, y) in tqdm(iterator):

    x = x.to(device)

    # Ensure only one value is unpacked
    output = model(x)
    if isinstance(output, tuple):
        y_pred, h = output
    else:
        y_pred = output
        h = None

    outputs.append(y_pred.cpu())
    intermediates.append(h.cpu() if h is not None else None)
    labels.append(y)

outputs = torch.cat(outputs, dim=0)
intermediates = torch.cat(intermediates, dim=0) if intermediates[0] is not None else None
labels = torch.cat(labels, dim=0)

return outputs, intermediates, labels

```

```
[ ]: outputs, intermediates, labels = get_representations(network,
                                                               trainloader,
                                                               device)
```

```

0%|          | 0/537 [00:00<?,
?it/s]/usr/lib/python3.10/multiprocessing/popen_fork.py:66: RuntimeWarning:
os.fork() was called. os.fork() is incompatible with multithreaded code, and JAX
is multithreaded, so this will likely lead to a deadlock.
    self.pid = os.fork()
100%|| 537/537 [03:06<00:00,  2.88it/s]
```

```
[ ]: def get_pca(data, n_components=2):
    pca = decomposition.PCA()
    pca.n_components = n_components
    pca_data = pca.fit_transform(data)
    return pca_data
```

```
[ ]: def plot_representations(data, labels, n_images=None):
    if n_images is not None:
        data = data[:n_images]
        labels = labels[:n_images]
    fig = plt.figure(figsize=(10, 10))
    ax = fig.add_subplot(111)
```

```
scatter = ax.scatter(data[:, 0], data[:, 1], c=labels, cmap='tab10')
handles, labels = scatter.legend_elements()
ax.legend(handles=handles, labels=labels)
```

```
[ ]: output_pca_data = get_pca(outputs)
```

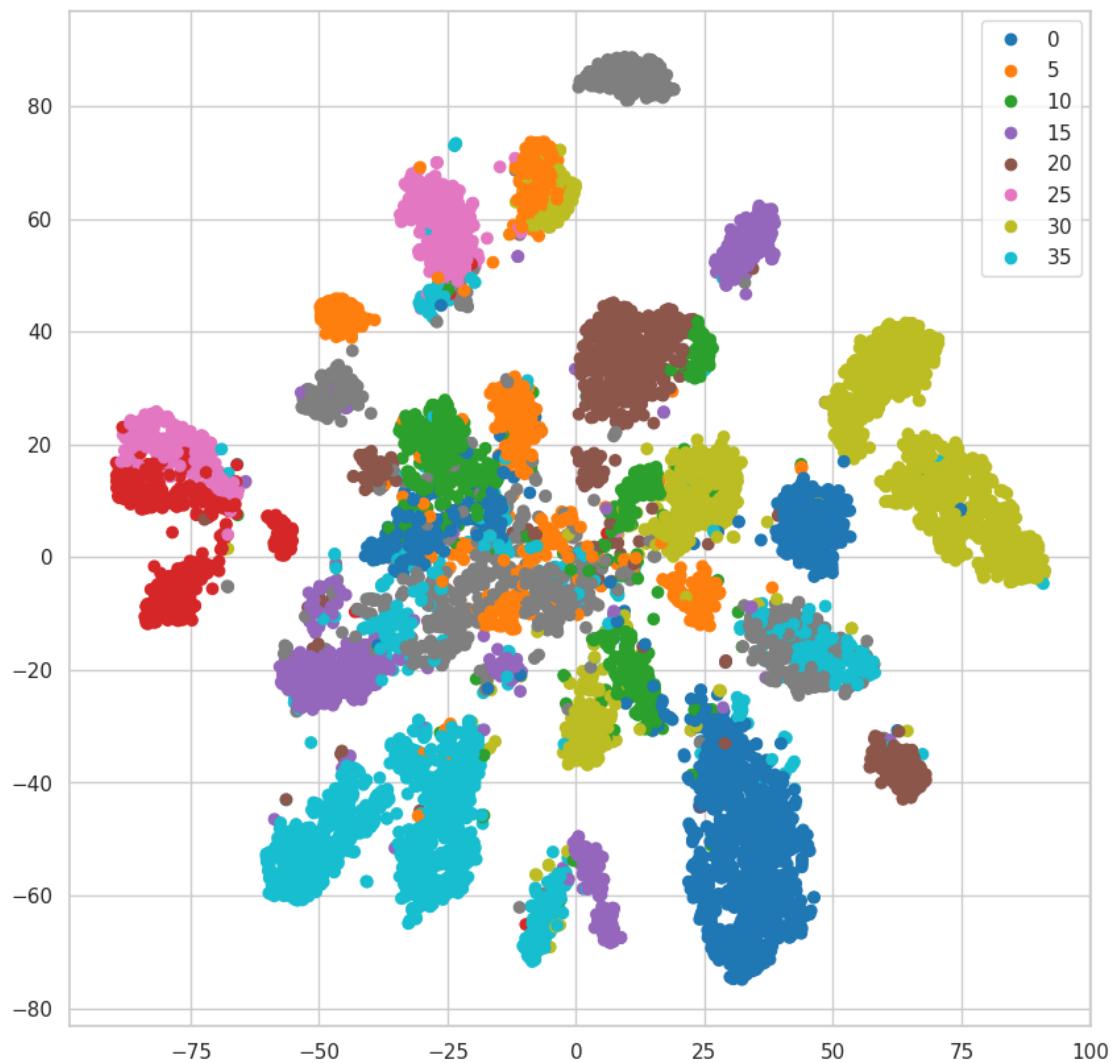
```
[ ]: plot_representations(output_pca_data, labels)
```



```
[ ]: def get_tsne(data, n_components=2, n_images=None):
    if n_images is not None:
        data = data[:n_images]
    tsne = manifold.TSNE(n_components=n_components, random_state=0)
    tsne_data = tsne.fit_transform(data)
    return tsne_data
```

```
[ ]: N_IMAGES = len(test_set)

output_tsne_data = get_tsne(outputs, n_images=N_IMAGES)
plot_representations(output_tsne_data, labels, n_images=N_IMAGES)
```



4 D. 2) Train the model over the grey images

```
[ ]: # Clone the repository
repo_url = "https://github.com/digitalepidemiologylab/
    ↳plantvillage_deeplearning_paper_dataset.git"
clone_dir = "plantvillage_deeplearning_paper_dataset"
subprocess.run(["git", "clone", repo_url, clone_dir])
os.chdir(clone_dir)
```

```

# Raw color images
classes = extract_folder(repo_url, clone_dir, "raw/grayscale")

Folder 'raw/grayscale' extracted successfully.
Directory 'Tomato___Late_blight' contains 1909 files.
Directory 'Apple___Cedar_apple_rust' contains 275 files.
Directory 'Corn_(maize)___Common_rust_' contains 1192 files.
Directory 'Peach___Bacterial_spot' contains 2297 files.
Directory 'Grape___Black_rot' contains 1180 files.
Directory 'Blueberry___healthy' contains 1502 files.
Directory 'Pepper,_bell___Bacterial_spot' contains 997 files.
Directory 'Grape___Leaf_blight_(Isariopsis_Leaf_Spot)' contains 1076 files.
Directory 'Peach___healthy' contains 360 files.
Directory 'Pepper,_bell___healthy' contains 1478 files.
Directory 'Corn_(maize)___healthy' contains 1162 files.
Directory 'Tomato___Spider_mites Two-spotted_spider_mite' contains 1676 files.
Directory 'Corn_(maize)___Northern_Leaf_Blight' contains 985 files.
Directory 'Orange___Haunglongbing_(Citrus_greening)' contains 5507 files.
Directory 'Squash___Powdery_mildew' contains 1835 files.
Directory 'Potato___healthy' contains 152 files.
Directory 'Apple___healthy' contains 1645 files.
Directory 'Strawberry___healthy' contains 456 files.
Directory 'Tomato___Septoria_leaf_spot' contains 1771 files.
Directory 'Tomato___Leaf_Mold' contains 952 files.
Directory 'Cherry_(including_sour)___healthy' contains 854 files.
Directory 'Tomato___Tomato_mosaic_virus' contains 373 files.
Directory 'Potato___Late_blight' contains 1000 files.
Directory 'Grape___healthy' contains 423 files.
Directory 'Potato___Early_blight' contains 1000 files.
Directory 'Apple___Black_rot' contains 621 files.
Directory 'Cherry_(including_sour)___Powdery_mildew' contains 1052 files.
Directory 'Corn_(maize)___Cercospora_leaf_spot_Gray_leaf_spot' contains 513
files.
Directory 'Tomato___healthy' contains 1591 files.
Directory 'Tomato___Target_Spot' contains 1404 files.
Directory 'Tomato___Tomato_Yellow_Leaf_Curl_Virus' contains 5357 files.
Directory 'Tomato___Bacterial_spot' contains 2127 files.
Directory 'Soybean___healthy' contains 5090 files.
Directory 'Apple___Apple_scab' contains 630 files.
Directory 'Raspberry___healthy' contains 371 files.
Directory 'Grape___Esca_(Black_Measles)' contains 1383 files.
Directory 'Strawberry___Leaf_scorch' contains 1109 files.
Directory 'Tomato___Early_blight' contains 1000 files.
Total number of files: 54305

total number of classes : '38'.
['Tomato___Late_blight', 'Apple___Cedar_apple_rust',

```

```
'Corn_(maize)___Common_rust_ ', 'Peach___Bacterial_spot', 'Grape___Black_rot',
'Blueberry___healthy', 'Pepper,_bell___Bacterial_spot',
'Grape___Leaf_blight_(Isariopsis_Leaf_Spot)', 'Peach___healthy',
'Pepper,_bell___healthy', 'Corn_(maize)___healthy', 'Tomato___Spider_mites Two-
spotted_spider_mite', 'Corn_(maize)___Northern_Leaf_Blight',
'Orange___Haunglongbing_(Citrus_greening)', 'Squash___Powdery_mildew',
'Potato___healthy', 'Apple___healthy', 'Strawberry___healthy',
'Tomato___Septoria_leaf_spot', 'Tomato___Leaf_Mold',
'Cherry_(including_sour)___healthy', 'Tomato___Tomato_mosaic_virus',
'Potato___Late_blight', 'Grape___healthy', 'Potato___Early_blight',
'Apple___Black_rot', 'Cherry_(including_sour)___Powdery_mildew',
'Corn_(maize)___Cercospora_leaf_spot_Gray_leaf_spot', 'Tomato___healthy',
'Tomato___Target_Spot', 'Tomato___Tomato_Yellow_Leaf_Curl_Virus',
'Tomato___Bacterial_spot', 'Soybean___healthy', 'Apple___Apple_scab',
'Raspberry___healthy', 'Grape___Esca_(Black_Measles)',
'Strawberry___Leaf_scorch', 'Tomato___Early_blight']
```

```
[ ]: # Splitting the dataset
extracted_folder = "raw/grayscale"
train_set = CropDiseaseDataset(root_dir=extracted_folder, train=True,
                                validation=False, gray_scale=True)
validation_set = CropDiseaseDataset(root_dir=extracted_folder, train=False,
                                    validation=True, gray_scale=True)
test_set = CropDiseaseDataset(root_dir=extracted_folder, train=False,
                               validation=False, gray_scale=True)

trainloader = torch.utils.data.DataLoader(train_set, batch_size=64,
                                         shuffle=True, num_workers=2)
validationloader = torch.utils.data.DataLoader(validation_set, batch_size=64,
                                               shuffle=False, num_workers=2)
testloader = torch.utils.data.DataLoader(test_set, batch_size=64, shuffle=False,
                                         num_workers=2)
```

```
[ ]: merged_dataset = ConcatDataset([train_set, validation_set, test_set])

# Create a single loader for the combined dataset
mergedloader = DataLoader(merged_dataset, batch_size=64, shuffle=True,
                           num_workers=2)
```

```
[ ]: class_counts = defaultdict(int)
for images, labels in mergedloader:
    for label in labels:
        class_counts[label.item()] += 1

class_labels = sorted(class_counts.keys())
counts = [class_counts[label] for label in class_labels]
colors = sns.color_palette("husl", len(class_labels))
```

```

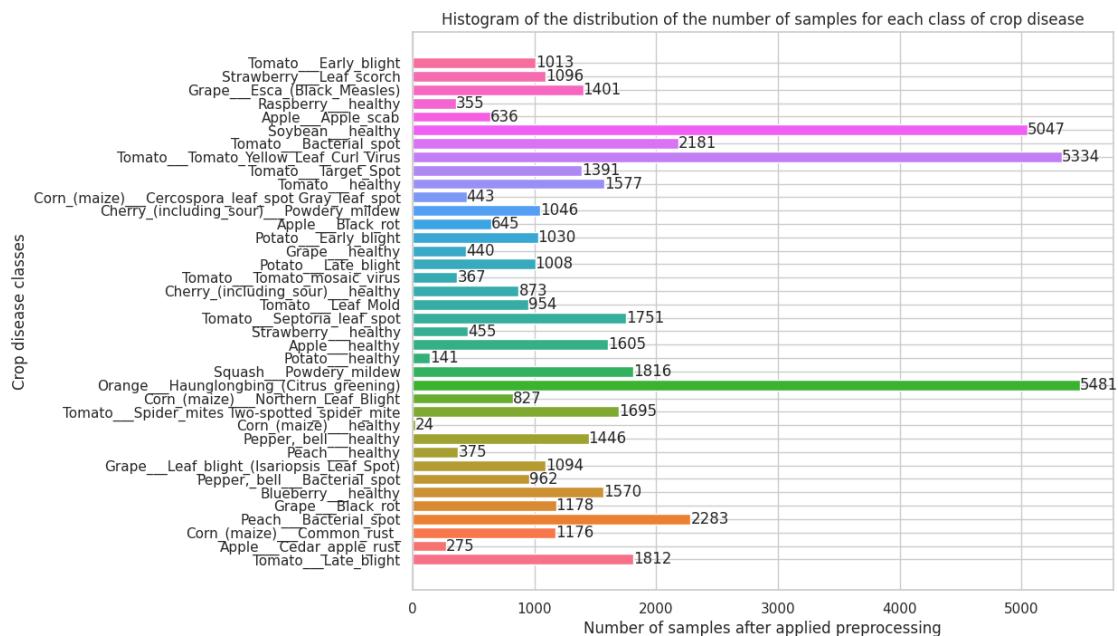
class_names = [train_set.classes[label] for label in class_labels]

sns.set(style="whitegrid")
plt.figure(figsize=(10, 8))
plt.barh(class_names, counts, color=colors)
plt.xlabel('Number of samples after applied preprocessing')
plt.ylabel('Crop disease classes')
plt.title('Histogram of the distribution of the number of samples for each class of crop disease')

for i, (count, name) in enumerate(zip(counts, class_names)):
    plt.text(count, i, str(count), ha='left', va='center')

plt.show()

```



```

[ ]: # Combining train and validation datasets for testing the model
combined_train_set = ConcatDataset([train_set, validation_set])
combined_train_loader = torch.utils.data.DataLoader(combined_train_set, batch_size=64, shuffle=True, num_workers=2)

```

```

[ ]: # Define hyperparameters
learning_rate = 1e-3
num_epochs = 25

network = CNN(grayscale=True).to(device)

```

```

network.apply(initialize_parameters)
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(network.parameters(), lr=learning_rate)

print(f'The model has {count_parameters(network)} trainable parameters')

```

The model has 6,846,182 trainable parameters

```
[ ]: # Train the model
train_avg_loss, validation_avg_loss, validation_accuracy, train_accuracy =
    ↪train(network, num_epochs, trainloader, validationloader, criterion,
    ↪optimizer, validation_phase=True)
```

```

Epoch [1/20], Validation Loss: 1.6054, Validation Accuracy: 54.95%, Train Loss:
2.6091, Train Accuracy: 36.80%
| Epoch Time: 0m 42s
Epoch [2/20], Validation Loss: 1.0772, Validation Accuracy: 67.71%, Train Loss:
1.3309, Train Accuracy: 61.05%
| Epoch Time: 0m 41s
Epoch [3/20], Validation Loss: 0.7704, Validation Accuracy: 76.12%, Train Loss:
0.9146, Train Accuracy: 71.57%
| Epoch Time: 0m 35s
Epoch [4/20], Validation Loss: 0.6666, Validation Accuracy: 78.38%, Train Loss:
0.7092, Train Accuracy: 77.35%
| Epoch Time: 0m 33s
Epoch [5/20], Validation Loss: 0.5722, Validation Accuracy: 81.86%, Train Loss:
0.5921, Train Accuracy: 81.03%
| Epoch Time: 0m 34s
Epoch [6/20], Validation Loss: 0.5278, Validation Accuracy: 83.27%, Train Loss:
0.5053, Train Accuracy: 83.48%
| Epoch Time: 0m 35s
Epoch [7/20], Validation Loss: 0.4972, Validation Accuracy: 83.91%, Train Loss:
0.4425, Train Accuracy: 85.19%
| Epoch Time: 0m 35s
Epoch [8/20], Validation Loss: 0.4520, Validation Accuracy: 85.43%, Train Loss:
0.4040, Train Accuracy: 86.79%
| Epoch Time: 0m 34s
Epoch [9/20], Validation Loss: 0.4134, Validation Accuracy: 86.72%, Train Loss:
0.3619, Train Accuracy: 88.12%
| Epoch Time: 0m 35s
Epoch [10/20], Validation Loss: 0.3941, Validation Accuracy: 87.16%, Train Loss:
0.3335, Train Accuracy: 89.00%
| Epoch Time: 0m 35s
Epoch [11/20], Validation Loss: 0.4001, Validation Accuracy: 87.50%, Train Loss:
0.3088, Train Accuracy: 89.83%
| Epoch Time: 0m 35s
Epoch [12/20], Validation Loss: 0.3507, Validation Accuracy: 88.84%, Train Loss:
0.2798, Train Accuracy: 90.71%

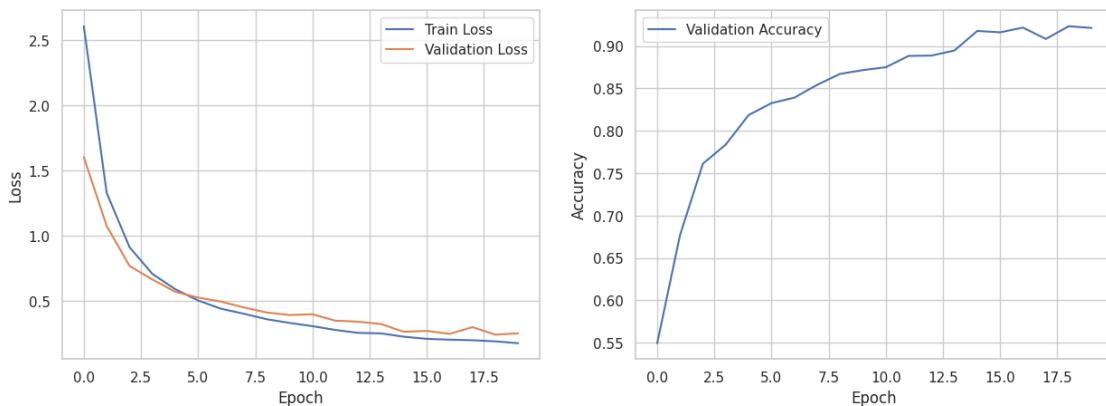
```

```

| Epoch Time: 0m 35s
Epoch [13/20], Validation Loss: 0.3432, Validation Accuracy: 88.88%, Train Loss: 0.2573, Train Accuracy: 91.49%
| Epoch Time: 0m 34s
Epoch [14/20], Validation Loss: 0.3248, Validation Accuracy: 89.47%, Train Loss: 0.2532, Train Accuracy: 91.45%
| Epoch Time: 0m 34s
Epoch [15/20], Validation Loss: 0.2655, Validation Accuracy: 91.78%, Train Loss: 0.2278, Train Accuracy: 92.47%
| Epoch Time: 0m 36s
Epoch [16/20], Validation Loss: 0.2724, Validation Accuracy: 91.62%, Train Loss: 0.2115, Train Accuracy: 93.07%
| Epoch Time: 0m 35s
Epoch [17/20], Validation Loss: 0.2497, Validation Accuracy: 92.17%, Train Loss: 0.2051, Train Accuracy: 93.28%
| Epoch Time: 0m 34s
Epoch [18/20], Validation Loss: 0.3011, Validation Accuracy: 90.83%, Train Loss: 0.2006, Train Accuracy: 93.28%
| Epoch Time: 0m 35s
Epoch [19/20], Validation Loss: 0.2435, Validation Accuracy: 92.34%, Train Loss: 0.1922, Train Accuracy: 93.76%
| Epoch Time: 0m 34s
Epoch [20/20], Validation Loss: 0.2541, Validation Accuracy: 92.13%, Train Loss: 0.1778, Train Accuracy: 94.07%
| Epoch Time: 0m 35s

```

```
[ ]: plot_training_results(train_avg_loss, validation_avg_loss, validation_accuracy, ↴is_validation=True)
```



```
[ ]: # Test the model
network = CNN(grayscale=True).to(device)
network.apply(initialize_parameters)
criterion = nn.CrossEntropyLoss()
```

```

optimizer = torch.optim.Adam(network.parameters(), lr=learning_rate)

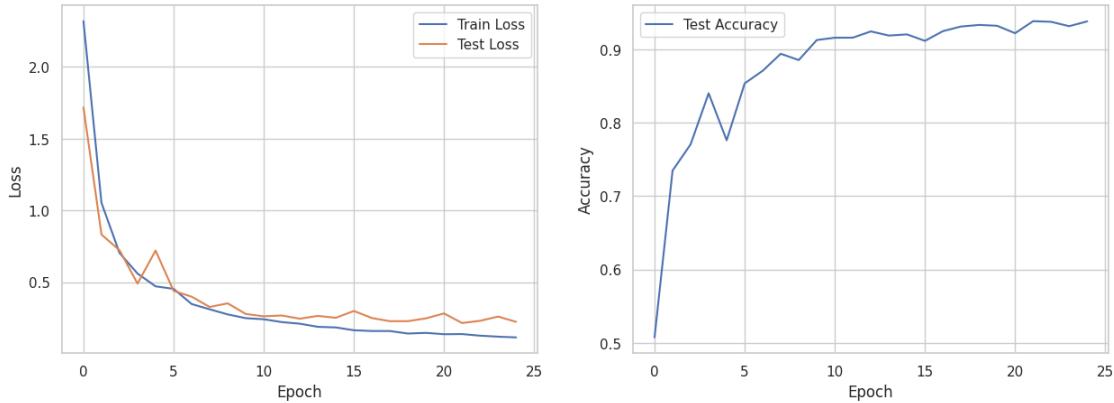
train_avg_loss, test_avg_loss, test_accuracy, train_accuracy = train(network, □
    ↵num_epochs, combined_train_loader, testloader, criterion, □
    ↵optimizer, validation_phase=False)

```

Epoch [1/25], Test Loss: 1.7172, Test Accuracy: 50.78%, Train Loss: 2.3190,
 Train Accuracy: 42.54%
 | Epoch Time: 0m 33s
 Epoch [2/25], Test Loss: 0.8323, Test Accuracy: 73.48%, Train Loss: 1.0529,
 Train Accuracy: 68.46%
 | Epoch Time: 0m 31s
 Epoch [3/25], Test Loss: 0.7229, Test Accuracy: 77.05%, Train Loss: 0.7042,
 Train Accuracy: 78.11%
 | Epoch Time: 0m 33s
 Epoch [4/25], Test Loss: 0.4907, Test Accuracy: 84.00%, Train Loss: 0.5606,
 Train Accuracy: 81.93%
 | Epoch Time: 0m 31s
 Epoch [5/25], Test Loss: 0.7209, Test Accuracy: 77.61%, Train Loss: 0.4717,
 Train Accuracy: 84.80%
 | Epoch Time: 0m 33s
 Epoch [6/25], Test Loss: 0.4400, Test Accuracy: 85.35%, Train Loss: 0.4541,
 Train Accuracy: 85.46%
 | Epoch Time: 0m 31s
 Epoch [7/25], Test Loss: 0.3997, Test Accuracy: 87.07%, Train Loss: 0.3484,
 Train Accuracy: 88.79%
 | Epoch Time: 0m 33s
 Epoch [8/25], Test Loss: 0.3278, Test Accuracy: 89.38%, Train Loss: 0.3116,
 Train Accuracy: 89.84%
 | Epoch Time: 0m 31s
 Epoch [9/25], Test Loss: 0.3531, Test Accuracy: 88.52%, Train Loss: 0.2761,
 Train Accuracy: 90.86%
 | Epoch Time: 0m 33s
 Epoch [10/25], Test Loss: 0.2792, Test Accuracy: 91.25%, Train Loss: 0.2498,
 Train Accuracy: 91.78%
 | Epoch Time: 0m 31s
 Epoch [11/25], Test Loss: 0.2628, Test Accuracy: 91.56%, Train Loss: 0.2423,
 Train Accuracy: 91.97%
 | Epoch Time: 0m 33s
 Epoch [12/25], Test Loss: 0.2685, Test Accuracy: 91.57%, Train Loss: 0.2226,
 Train Accuracy: 92.61%
 | Epoch Time: 0m 31s
 Epoch [13/25], Test Loss: 0.2460, Test Accuracy: 92.43%, Train Loss: 0.2113,
 Train Accuracy: 93.10%
 | Epoch Time: 0m 32s
 Epoch [14/25], Test Loss: 0.2656, Test Accuracy: 91.86%, Train Loss: 0.1896,
 Train Accuracy: 93.83%
 | Epoch Time: 0m 32s

```
Epoch [15/25], Test Loss: 0.2523, Test Accuracy: 92.02%, Train Loss: 0.1852,  
Train Accuracy: 93.91%  
| Epoch Time: 0m 32s  
Epoch [16/25], Test Loss: 0.3001, Test Accuracy: 91.14%, Train Loss: 0.1656,  
Train Accuracy: 94.51%  
| Epoch Time: 0m 33s  
Epoch [17/25], Test Loss: 0.2508, Test Accuracy: 92.46%, Train Loss: 0.1604,  
Train Accuracy: 94.67%  
| Epoch Time: 0m 31s  
Epoch [18/25], Test Loss: 0.2286, Test Accuracy: 93.09%, Train Loss: 0.1603,  
Train Accuracy: 94.84%  
| Epoch Time: 0m 33s  
Epoch [19/25], Test Loss: 0.2289, Test Accuracy: 93.31%, Train Loss: 0.1426,  
Train Accuracy: 95.45%  
| Epoch Time: 0m 32s  
Epoch [20/25], Test Loss: 0.2485, Test Accuracy: 93.19%, Train Loss: 0.1473,  
Train Accuracy: 95.37%  
| Epoch Time: 0m 33s  
Epoch [21/25], Test Loss: 0.2828, Test Accuracy: 92.18%, Train Loss: 0.1380,  
Train Accuracy: 95.54%  
| Epoch Time: 0m 31s  
Epoch [22/25], Test Loss: 0.2161, Test Accuracy: 93.82%, Train Loss: 0.1390,  
Train Accuracy: 95.63%  
| Epoch Time: 0m 32s  
Epoch [23/25], Test Loss: 0.2313, Test Accuracy: 93.73%, Train Loss: 0.1278,  
Train Accuracy: 95.90%  
| Epoch Time: 0m 32s  
Epoch [24/25], Test Loss: 0.2608, Test Accuracy: 93.14%, Train Loss: 0.1208,  
Train Accuracy: 96.08%  
| Epoch Time: 0m 32s  
Epoch [25/25], Test Loss: 0.2242, Test Accuracy: 93.79%, Train Loss: 0.1158,  
Train Accuracy: 96.30%  
| Epoch Time: 0m 31s
```

```
[ ]: plot_training_results(train_avg_loss, test_avg_loss, test_accuracy, ▶  
    ↪is_validation=False)
```



Examining the model.

```
[ ]: network.load_state_dict(torch.load('model.pt'))
```

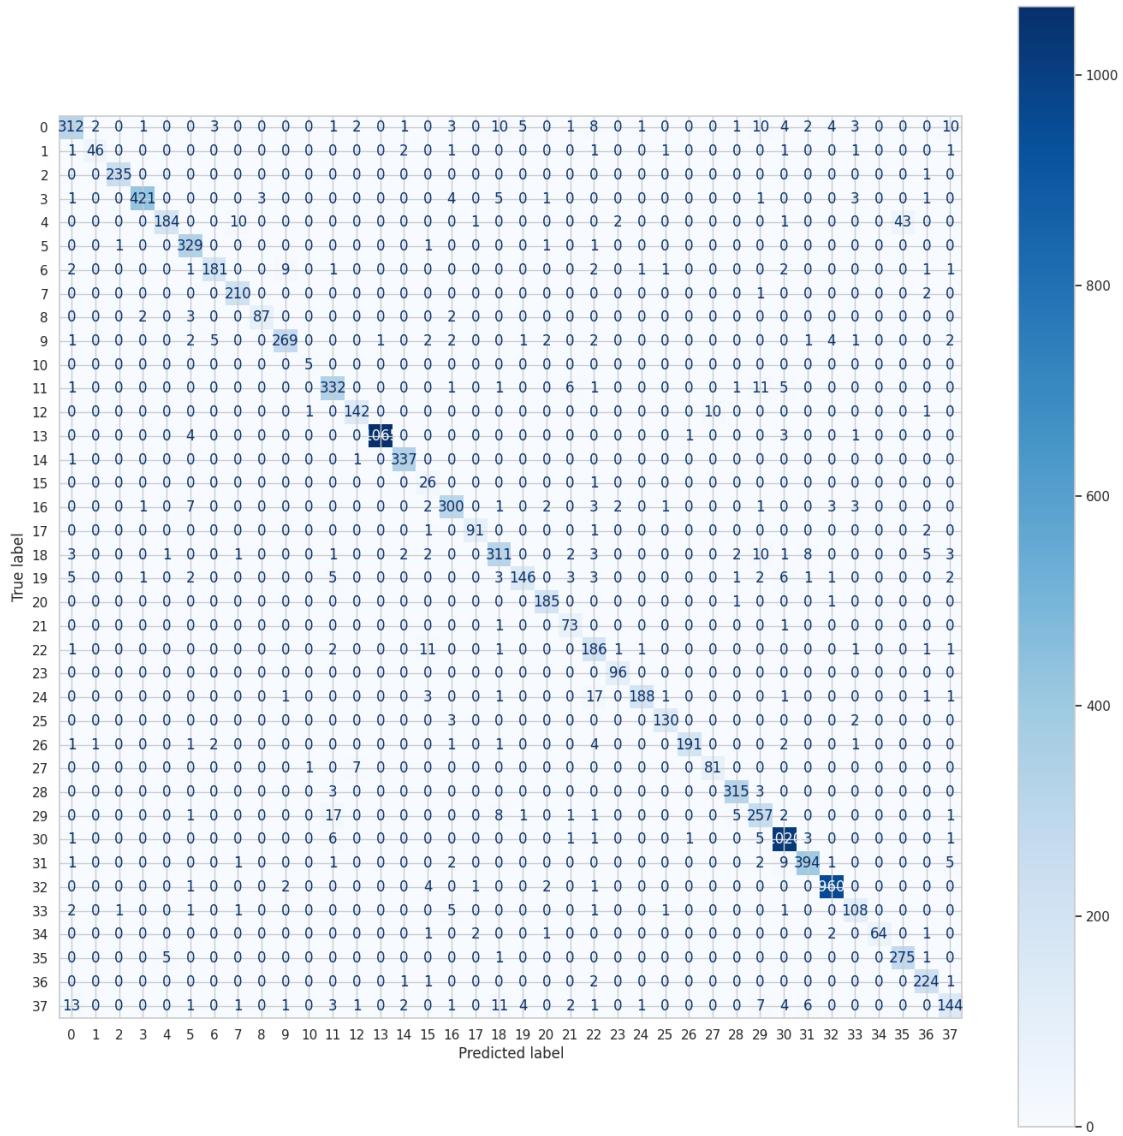
```
[ ]: <All keys matched successfully>
```

```
[ ]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
images, labels, probs, corrects = get_predictions(network, testloader, device)
pred_labels = torch.argmax(probs, 1)

print(f"There are {len(corrects)} correct predictions.")
```

There are 10562 correct predictions.

```
[ ]: plot_confusion_matrix(labels, pred_labels)
```



```
[ ]: # Now, extract and sort incorrect examples
incorrect_examples = []
for image, label, prob, correct in zip(images, labels, probs, corrects):
    if not correct:
        incorrect_examples.append((image, label, prob))

incorrect_examples.sort(reverse=True, key=lambda x: torch.max(x[2], dim=0).
                           values)
print(f"There are {len(incorrect_examples)} incorrect predictions.")
```

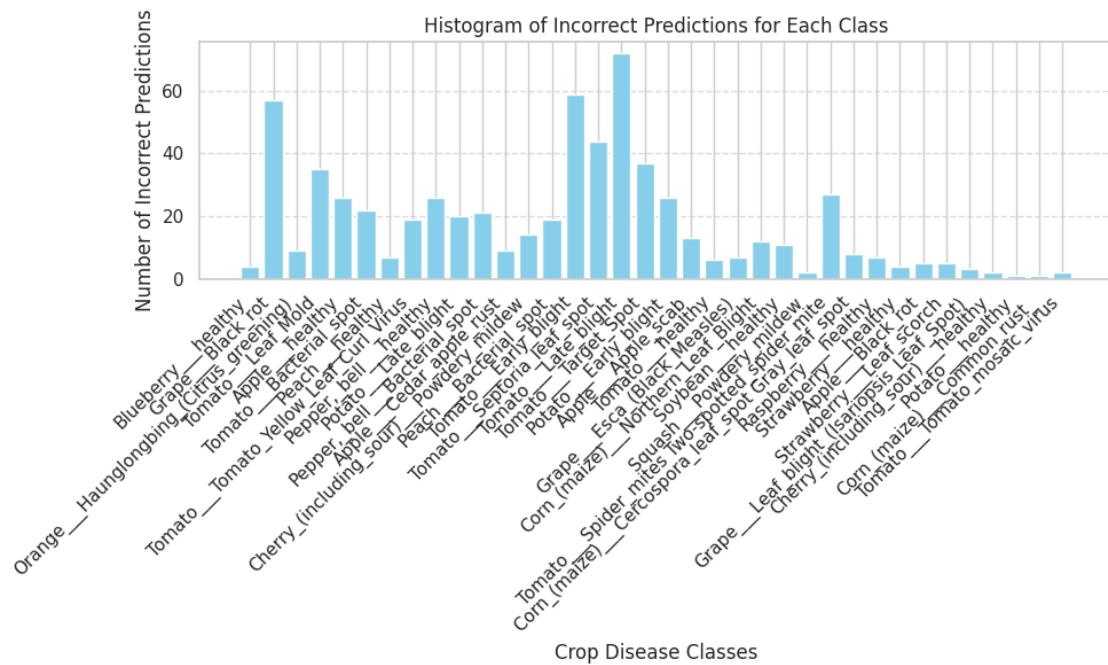
There are 642 incorrect predictions.

```
[ ]: incorrect_counts = {}

# Count the occurrences of each class label in the incorrect examples
for _, label, _ in incorrect_examples:
    if label.item() not in incorrect_counts:
        incorrect_counts[label.item()] = 1
    else:
        incorrect_counts[label.item()] += 1

class_names = [train_set.classes[label] for label in incorrect_counts.keys()]

# Plotting the histogram
plt.figure(figsize=(10, 6))
plt.bar(class_names, incorrect_counts.values(), color='skyblue')
plt.xlabel('Crop Disease Classes')
plt.ylabel('Number of Incorrect Predictions')
plt.title('Histogram of Incorrect Predictions for Each Class')
plt.xticks(rotation=45, ha='right')
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()
```



```
[ ]: data = {'Class Name': class_names, 'Incorrect Counts': list(incorrect_counts.  
    ↵values())}  
df1 = pd.DataFrame(data)
```

```
df1.head(5)
```

```
[ ]:          Class Name  Incorrect Counts
0           Blueberry___healthy               4
1            Grape___Black_rot              57
2 Orange___Haunglongbing_(Citrus_greening)      9
3           Tomato___Leaf_Mold              35
4           Apple___healthy                26
```

```
[ ]: class_counts = defaultdict(int)
for images, labels in testloader:
    for label in labels:
        class_counts[label.item()] += 1

class_labels = sorted(class_counts.keys())
counts = [class_counts[label] for label in class_labels]
colors = sns.color_palette("husl", len(class_labels))

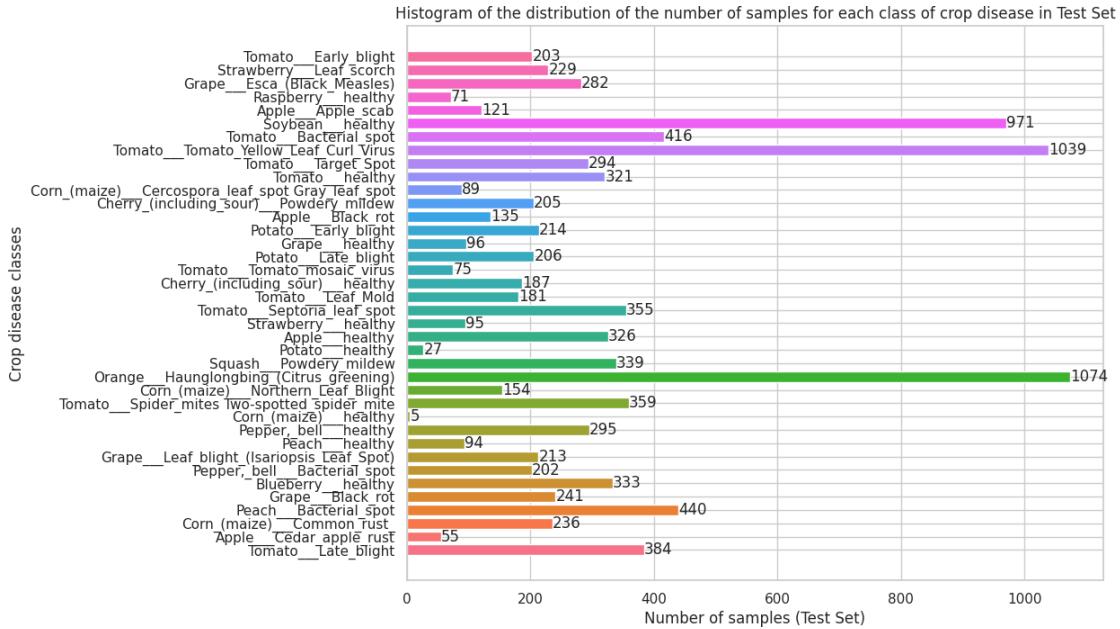
class_names = [train_set.classes[label] for label in class_labels]

sns.set(style="whitegrid")
plt.figure(figsize=(10, 8))
plt.barh(class_names, counts, color=colors)
plt.xlabel('Number of samples (Test Set)')
plt.ylabel('Crop disease classes')
plt.title('Histogram of the distribution of the number of samples for each class\u2014of crop disease in Test Set')

for i, (count, name) in enumerate(zip(counts, class_names)):
    plt.text(count, i, str(count), ha='left', va='center')

plt.show()
```

```
/usr/lib/python3.10/multiprocessing/popen_fork.py:66: RuntimeWarning: os.fork()
was called. os.fork() is incompatible with multithreaded code, and JAX is
multithreaded, so this will likely lead to a deadlock.
    self.pid = os.fork()
/usr/lib/python3.10/multiprocessing/popen_fork.py:66: RuntimeWarning: os.fork()
was called. os.fork() is incompatible with multithreaded code, and JAX is
multithreaded, so this will likely lead to a deadlock.
    self.pid = os.fork()
```



```
[ ]: data = {'Class Name': class_names, 'Counts': counts}
df2 = pd.DataFrame(data)
df2.head(5)
```

```
[ ]:      Class Name  Counts
0      Tomato_Late_blight    384
1      Apple_Cedar_apple_rust     55
2  Corn_(maize)_Common_rust_    236
3      Peach_Bacterial_spot    440
4      Grape_Black_rot        241
```

```
[ ]: merged_df = pd.merge(df1, df2, on='Class Name', suffixes=('_incorrect', '_test'))
merged_df.head(3)
```

```
[ ]:      Class Name  Incorrect Counts  Counts
0      Blueberry_healthy           4       333
1      Grape_Black_rot          57       241
2  Orange_Haunglongbing_(Citrus_greening)         9       1074
```

```
[ ]: # Calculate the rate of success for each class in percentage
merged_df['Success Rate (%)'] = ((merged_df['Counts'] - merged_df['Incorrect_Counts']) / merged_df['Counts']) * 100
merged_df
```

```
[ ]:      Class Name  Incorrect Counts \
0      Blueberry_healthy           4
```

1		Grape___Black_rot	57
2	Orange___Haunglongbing_(Citrus_greening)		9
3		Tomato___Leaf_Mold	35
4		Apple___healthy	26
5		Tomato___Bacterial_spot	22
6		Peach___healthy	7
7	Tomato___Tomato_Yellow_Leaf_Curl_Virus		19
8		Pepper,_bell___healthy	26
9		Potato___Late_blight	20
10		Pepper,_bell___Bacterial_spot	21
11		Apple___Cedar_apple_rust	9
12	Cherry_(including_sour)___Powdery_mildew		14
13		Peach___Bacterial_spot	19
14		Tomato___Early_blight	59
15		Tomato___Septoria_leaf_spot	44
16		Tomato___Late_blight	72
17		Tomato___Target_Spot	37
18		Potato___Early_blight	26
19		Apple___Apple_scab	13
20		Tomato___healthy	6
21		Grape___Esca_(Black_Measles)	7
22	Corn_(maize)___Northern_Leaf_Blight		12
23		Soybean___healthy	11
24		Squash___Powdery_mildew	2
25	Tomato___Spider_mites Two-spotted_spider_mite		27
26	Corn_(maize)___Cercospora_leaf_spot_Gray_leaf...		8
27		Raspberry___healthy	7
28		Strawberry___healthy	4
29		Apple___Black_rot	5
30		Strawberry___Leaf_scorch	5
31	Grape___Leaf_blight_(Isariopsis_Leaf_Spot)		3
32		Cherry_(including_sour)___healthy	2
33		Potato___healthy	1
34		Corn_(maize)___Common_rust_	1
35		Tomato___Tomato_mosaic_virus	2

	Counts	Success Rate (%)
0	333	98.798799
1	241	76.348548
2	1074	99.162011
3	181	80.662983
4	326	92.024540
5	416	94.711538
6	94	92.553191
7	1039	98.171319
8	295	91.186441
9	206	90.291262

10	202	89.603960
11	55	83.636364
12	205	93.170732
13	440	95.681818
14	203	70.935961
15	355	87.605634
16	384	81.250000
17	294	87.414966
18	214	87.850467
19	121	89.256198
20	321	98.130841
21	282	97.517730
22	154	92.207792
23	971	98.867147
24	339	99.410029
25	359	92.479109
26	89	91.011236
27	71	90.140845
28	95	95.789474
29	135	96.296296
30	229	97.816594
31	213	98.591549
32	187	98.930481
33	27	96.296296
34	236	99.576271
35	75	97.333333

```
[ ]: N_IMAGES = 25
plot_most_incorrect(incorrect_examples, N_IMAGES)
```



[]:

[]:

[]:

5 D. 3) Train the model over the segmented images

```
[4]: repo_url = "https://github.com/digitalepidemiologylab/
    ↪plantvillage_deeplearning_paper_dataset.git"
clone_dir = "plantvillage_deeplearning_paper_dataset"
subprocess.run(["git", "clone", repo_url, clone_dir])
os.chdir(clone_dir)

# Raw segmented images
classes = extract_folder(repo_url, clone_dir, "raw/segmented")
```

Folder 'raw/segmented' extracted successfully.
 Directory 'Apple___Black_rot' contains 621 files.
 Directory 'Tomato___Leaf_Mold' contains 952 files.
 Directory 'Raspberry___healthy' contains 371 files.
 Directory 'Tomato___Tomato_Yellow_Leaf_Curl_Virus' contains 5357 files.
 Directory 'Apple___Apple_scab' contains 630 files.
 Directory 'Tomato___Early_blight' contains 1000 files.
 Directory 'Strawberry___healthy' contains 456 files.
 Directory 'Apple___Cedar_apple_rust' contains 275 files.

```
Directory 'Soybean___healthy' contains 5090 files.
Directory 'Potato___Late_blight' contains 1000 files.
Directory 'Peach___Bacterial_spot' contains 2297 files.
Directory 'Pepper,_bell___healthy' contains 1478 files.
Directory 'Grape___Leaf_blight_(Isariopsis_Leaf_Spot)' contains 1076 files.
Directory 'Strawberry___Leaf_scorch' contains 1109 files.
Directory 'Tomato___Tomato_mosaic_virus' contains 373 files.
Directory 'Grape___healthy' contains 423 files.
Directory 'Tomato___Target_Spot' contains 1404 files.
Directory 'Squash___Powdery_mildew' contains 1835 files.
Directory 'Corn_(maize)___Cercospora_leaf_spot_Gray_leaf_spot' contains 513
files.
Directory 'Blueberry___healthy' contains 1502 files.
Directory 'Tomato___Spider_mites Two-spotted_spider_mite' contains 1676 files.
Directory 'Peach___healthy' contains 360 files.
Directory 'Grape___Esca_(Black_Measles)' contains 1384 files.
Directory 'Corn_(maize)___Common_rust_' contains 1192 files.
Directory 'Corn_(maize)___Northern_Leaf_Blight' contains 985 files.
Directory 'Apple___healthy' contains 1645 files.
Directory 'Tomato___Septoria_leaf_spot' contains 1771 files.
Directory 'Orange___Haunglongbing_(Citrus_greening)' contains 5507 files.
Directory 'Cherry_(including_sour)___healthy' contains 854 files.
Directory 'Tomato___Bacterial_spot' contains 2127 files.
Directory 'Potato___Early_blight' contains 1000 files.
Directory 'Corn_(maize)___healthy' contains 1162 files.
Directory 'Cherry_(including_sour)___Powdery_mildew' contains 1052 files.
Directory 'Tomato___healthy' contains 1591 files.
Directory 'Potato___healthy' contains 152 files.
Directory 'Pepper,_bell___Bacterial_spot' contains 997 files.
Directory 'Tomato___Late_blight' contains 1909 files.
Directory 'Grape___Black_rot' contains 1180 files.
Total number of files: 54306
```

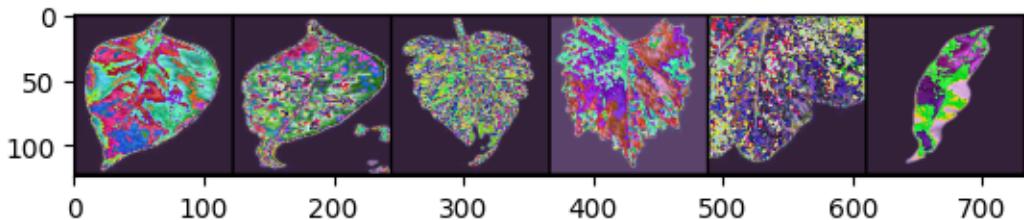
```
total number of classes : '38'.
['Apple___Black_rot', 'Tomato___Leaf_Mold', 'Raspberry___healthy',
'Tomato___Tomato_Yellow_Leaf_Curl_Virus', 'Apple___Apple_scab',
'Tomato___Early_blight', 'Strawberry___healthy', 'Apple___Cedar_apple_rust',
'Soybean___healthy', 'Potato___Late_blight', 'Peach___Bacterial_spot',
'Pepper,_bell___healthy', 'Grape___Leaf_blight_(Isariopsis_Leaf_Spot)',
'Strawberry___Leaf_scorch', 'Tomato___Tomato_mosaic_virus', 'Grape___healthy',
'Tomato___Target_Spot', 'Squash___Powdery_mildew',
'Corn_(maize)___Cercospora_leaf_spot_Gray_leaf_spot', 'Blueberry___healthy',
'Tomato___Spider_mites Two-spotted_spider_mite', 'Peach___healthy',
'Grape___Esca_(Black_Measles)', 'Corn_(maize)___Common_rust_',
'Corn_(maize)___Northern_Leaf_Blight', 'Apple___healthy',
'Tomato___Septoria_leaf_spot', 'Orange___Haunglongbing_(Citrus_greening)',
'Cherry_(including_sour)___healthy', 'Tomato___Bacterial_spot',
'Potato___Early_blight', 'Corn_(maize)___healthy',
```

```
'Cherry_(including_sour)___Powdery_mildew', 'Tomato___healthy',
'Potato___healthy', 'Pepper,_bell___Bacterial_spot', 'Tomato___Late_blight',
'Grape___Black_rot']
```

```
[ ]: extracted_folder = "raw/segmented"
my_dataset = CropDiseaseDataset(root_dir=extracted_folder, train=True, □
    ↪gray_scale=False, segmented=True)
train_loader = DataLoader(my_dataset, batch_size=6, shuffle=True, num_workers=2)

def show_images(img):
    plt.imshow(transforms.functional.to_pil_image(img))
    plt.show()

images, labels = next(iter(train_loader))
show_images(utils.make_grid(images))
print(*[my_dataset.classes[l] for l in labels])
```



Pepper,_bell___healthy Cherry_(including_sour)___Powdery_mildew Tomato___healthy
Grape___Black_rot Squash___Powdery_mildew Peach___Bacterial_spot

```
[11]: # Define hyperparameters
learning_rate = 1e-3
num_epochs = 25

network = CNN(grayscale=False).to(device)
network.apply(initialize_parameters)
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(network.parameters(), lr=learning_rate)

print(f'The model has {count_parameters(network):,} trainable parameters')
```

The model has 6,846,758 trainable parameters

```
[12]: # Splitting the dataset
extracted_folder = "raw/segmented"
train_set = CropDiseaseDataset(root_dir=extracted_folder, train=True, □
    ↪validation=False, gray_scale=False, segmented=True)
```

```

validation_set = CropDiseaseDataset(root_dir=extracted_folder, train=False,
                                   validation=True, gray_scale=False, segmented=True)
test_set = CropDiseaseDataset(root_dir=extracted_folder, train=False,
                               validation=False, gray_scale=False, segmented=True)

trainloader = torch.utils.data.DataLoader(train_set, batch_size=64,
                                         shuffle=True, num_workers=2)
validationloader = torch.utils.data.DataLoader(validation_set, batch_size=64,
                                               shuffle=False, num_workers=2)
testloader = torch.utils.data.DataLoader(test_set, batch_size=64, shuffle=False,
                                         num_workers=2)

```

[13]: merged_dataset = ConcatDataset([train_set, validation_set, test_set])

```

# Create a single loader for the combined dataset
mergedloader = DataLoader(merged_dataset, batch_size=64, shuffle=True,
                           num_workers=2)

```

[]:

```

class_counts = defaultdict(int)
for images, labels in mergedloader:
    for label in labels:
        class_counts[label.item()] += 1

class_labels = sorted(class_counts.keys())
counts = [class_counts[label] for label in class_labels]
colors = sns.color_palette("husl", len(class_labels))

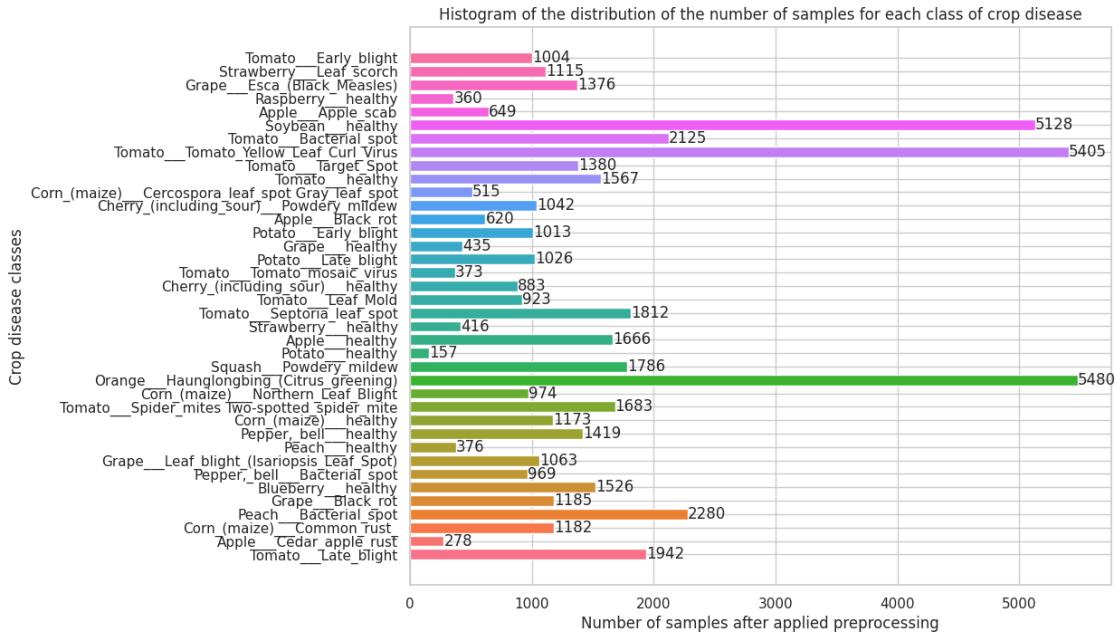
class_names = [train_set.classes[label] for label in class_labels]

sns.set(style="whitegrid")
plt.figure(figsize=(10, 8))
plt.barh(class_names, counts, color=colors)
plt.xlabel('Number of samples after applied preprocessing')
plt.ylabel('Crop disease classes')
plt.title('Histogram of the distribution of the number of samples for each class  
of crop disease')

for i, (count, name) in enumerate(zip(counts, class_names)):
    plt.text(count, i, str(count), ha='left', va='center')

plt.show()

```

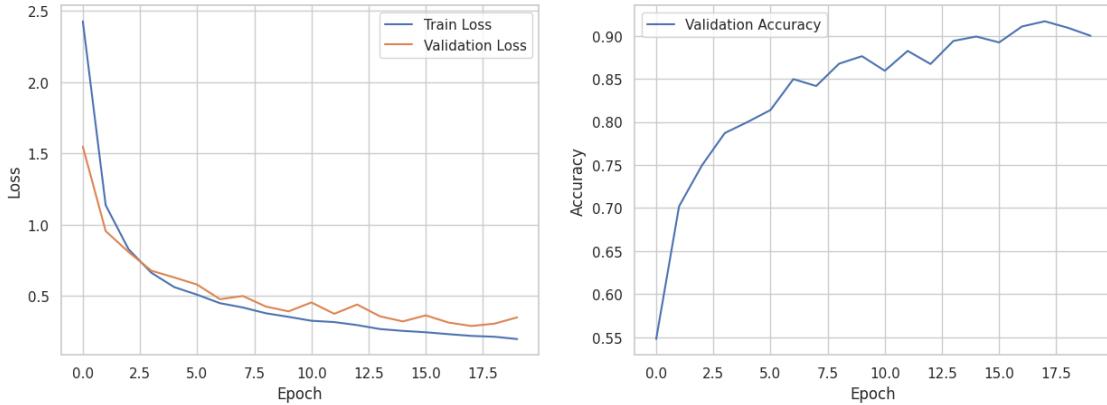


```
[ ]: # Train the models
train_avg_loss, validation_avg_loss, validation_accuracy, train_accuracy =
    →train(network, num_epochs, trainloader, validationloader, criterion,
    →optimizer, validation_phase=True)

Epoch [1/20], Validation Loss: 1.5477, Validation Accuracy: 54.78%, Train Loss:
2.4259, Train Accuracy: 40.44%
| Epoch Time: 0m 33s
Epoch [2/20], Validation Loss: 0.9554, Validation Accuracy: 70.19%, Train Loss:
1.1370, Train Accuracy: 65.83%
| Epoch Time: 0m 31s
Epoch [3/20], Validation Loss: 0.8088, Validation Accuracy: 74.94%, Train Loss:
0.8282, Train Accuracy: 73.92%
| Epoch Time: 0m 31s
Epoch [4/20], Validation Loss: 0.6777, Validation Accuracy: 78.70%, Train Loss:
0.6631, Train Accuracy: 78.95%
| Epoch Time: 0m 31s
Epoch [5/20], Validation Loss: 0.6303, Validation Accuracy: 79.99%, Train Loss:
0.5627, Train Accuracy: 81.72%
| Epoch Time: 0m 32s
Epoch [6/20], Validation Loss: 0.5804, Validation Accuracy: 81.38%, Train Loss:
0.5094, Train Accuracy: 83.40%
| Epoch Time: 0m 33s
Epoch [7/20], Validation Loss: 0.4780, Validation Accuracy: 84.96%, Train Loss:
0.4500, Train Accuracy: 85.12%
| Epoch Time: 0m 31s
Epoch [8/20], Validation Loss: 0.5004, Validation Accuracy: 84.16%, Train Loss:
```

```
0.4199, Train Accuracy: 86.27%
| Epoch Time: 0m 31s
Epoch [9/20], Validation Loss: 0.4268, Validation Accuracy: 86.75%, Train Loss:
0.3796, Train Accuracy: 87.57%
| Epoch Time: 0m 33s
Epoch [10/20], Validation Loss: 0.3928, Validation Accuracy: 87.62%, Train Loss:
0.3542, Train Accuracy: 88.37%
| Epoch Time: 0m 31s
Epoch [11/20], Validation Loss: 0.4552, Validation Accuracy: 85.93%, Train Loss:
0.3269, Train Accuracy: 89.06%
| Epoch Time: 0m 32s
Epoch [12/20], Validation Loss: 0.3755, Validation Accuracy: 88.24%, Train Loss:
0.3175, Train Accuracy: 89.73%
| Epoch Time: 0m 32s
Epoch [13/20], Validation Loss: 0.4408, Validation Accuracy: 86.72%, Train Loss:
0.2964, Train Accuracy: 90.19%
| Epoch Time: 0m 32s
Epoch [14/20], Validation Loss: 0.3577, Validation Accuracy: 89.39%, Train Loss:
0.2687, Train Accuracy: 91.12%
| Epoch Time: 0m 31s
Epoch [15/20], Validation Loss: 0.3222, Validation Accuracy: 89.91%, Train Loss:
0.2559, Train Accuracy: 91.79%
| Epoch Time: 0m 31s
Epoch [16/20], Validation Loss: 0.3646, Validation Accuracy: 89.22%, Train Loss:
0.2461, Train Accuracy: 92.11%
| Epoch Time: 0m 31s
Epoch [17/20], Validation Loss: 0.3141, Validation Accuracy: 91.07%, Train Loss:
0.2329, Train Accuracy: 92.29%
| Epoch Time: 0m 31s
Epoch [18/20], Validation Loss: 0.2903, Validation Accuracy: 91.68%, Train Loss:
0.2206, Train Accuracy: 92.67%
| Epoch Time: 0m 31s
Epoch [19/20], Validation Loss: 0.3063, Validation Accuracy: 90.93%, Train Loss:
0.2148, Train Accuracy: 92.97%
| Epoch Time: 0m 31s
Epoch [20/20], Validation Loss: 0.3506, Validation Accuracy: 89.99%, Train Loss:
0.1986, Train Accuracy: 93.47%
| Epoch Time: 0m 33s
```

```
[ ]: plot_training_results(train_avg_loss, validation_avg_loss, validation_accuracy,  
                         ↪is_validation=True)
```



```
[14]: # Combining train and validation datasets for testing the model
combined_train_set = ConcatDataset([train_set, validation_set])
combined_train_loader = torch.utils.data.DataLoader(combined_train_set, □
    ↪batch_size=64, shuffle=True, num_workers=2)

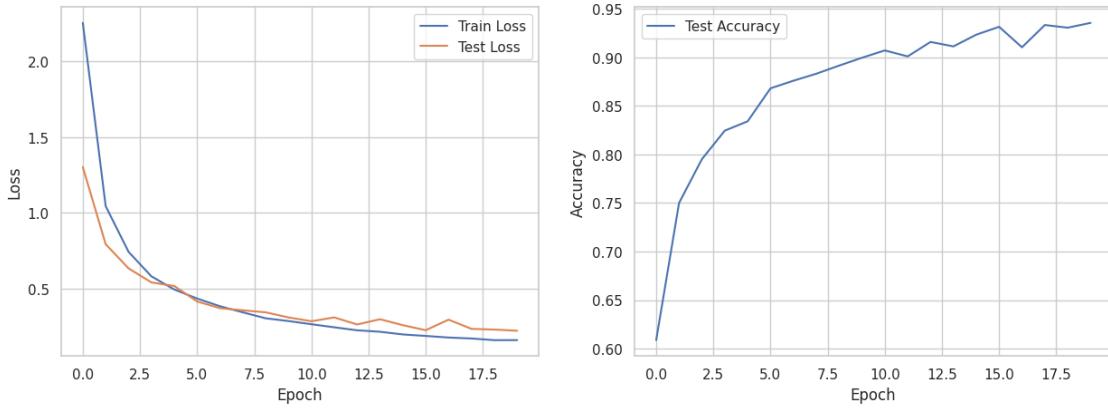
[20]: # Test the model
num_epochs = 20
network = CNN(grayscale=False).to(device)
network.apply(initialize_parameters)
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(network.parameters(), lr=learning_rate)

train_avg_loss, test_avg_loss, test_accuracy, train_accuracy = train(network, □
    ↪num_epochs, combined_train_loader, testloader, criterion, □
    ↪optimizer, validation_phase=False)
```

Epoch [1/20], Test Loss: 1.3019, Test Accuracy: 60.88%, Train Loss: 2.2548, Train Accuracy: 43.15%
| Epoch Time: 0m 46s
Epoch [2/20], Test Loss: 0.7924, Test Accuracy: 75.00%, Train Loss: 1.0440, Train Accuracy: 68.17%
| Epoch Time: 0m 46s
Epoch [3/20], Test Loss: 0.6328, Test Accuracy: 79.51%, Train Loss: 0.7420, Train Accuracy: 76.63%
| Epoch Time: 0m 42s
Epoch [4/20], Test Loss: 0.5399, Test Accuracy: 82.44%, Train Loss: 0.5800, Train Accuracy: 81.23%
| Epoch Time: 0m 42s
Epoch [5/20], Test Loss: 0.5161, Test Accuracy: 83.40%, Train Loss: 0.4931, Train Accuracy: 84.05%
| Epoch Time: 0m 43s
Epoch [6/20], Test Loss: 0.4137, Test Accuracy: 86.80%, Train Loss: 0.4333, Train Accuracy: 85.87%

```
| Epoch Time: 0m 42s
Epoch [7/20], Test Loss: 0.3694, Test Accuracy: 87.58%, Train Loss: 0.3824,
Train Accuracy: 87.54%
| Epoch Time: 0m 43s
Epoch [8/20], Test Loss: 0.3554, Test Accuracy: 88.30%, Train Loss: 0.3420,
Train Accuracy: 88.88%
| Epoch Time: 0m 42s
Epoch [9/20], Test Loss: 0.3423, Test Accuracy: 89.14%, Train Loss: 0.3026,
Train Accuracy: 90.14%
| Epoch Time: 0m 43s
Epoch [10/20], Test Loss: 0.3081, Test Accuracy: 89.95%, Train Loss: 0.2845,
Train Accuracy: 90.70%
| Epoch Time: 0m 44s
Epoch [11/20], Test Loss: 0.2834, Test Accuracy: 90.70%, Train Loss: 0.2635,
Train Accuracy: 91.45%
| Epoch Time: 0m 44s
Epoch [12/20], Test Loss: 0.3084, Test Accuracy: 90.09%, Train Loss: 0.2432,
Train Accuracy: 92.20%
| Epoch Time: 0m 42s
Epoch [13/20], Test Loss: 0.2618, Test Accuracy: 91.58%, Train Loss: 0.2233,
Train Accuracy: 92.75%
| Epoch Time: 0m 43s
Epoch [14/20], Test Loss: 0.2966, Test Accuracy: 91.12%, Train Loss: 0.2138,
Train Accuracy: 93.09%
| Epoch Time: 0m 43s
Epoch [15/20], Test Loss: 0.2572, Test Accuracy: 92.32%, Train Loss: 0.1959,
Train Accuracy: 93.64%
| Epoch Time: 0m 43s
Epoch [16/20], Test Loss: 0.2239, Test Accuracy: 93.13%, Train Loss: 0.1861,
Train Accuracy: 94.06%
| Epoch Time: 0m 43s
Epoch [17/20], Test Loss: 0.2936, Test Accuracy: 91.02%, Train Loss: 0.1756,
Train Accuracy: 94.41%
| Epoch Time: 0m 44s
Epoch [18/20], Test Loss: 0.2328, Test Accuracy: 93.32%, Train Loss: 0.1692,
Train Accuracy: 94.54%
| Epoch Time: 0m 43s
Epoch [19/20], Test Loss: 0.2282, Test Accuracy: 93.04%, Train Loss: 0.1582,
Train Accuracy: 94.95%
| Epoch Time: 0m 44s
Epoch [20/20], Test Loss: 0.2204, Test Accuracy: 93.54%, Train Loss: 0.1586,
Train Accuracy: 94.99%
| Epoch Time: 0m 43s
```

```
[21]: plot_training_results(train_avg_loss, test_avg_loss, test_accuracy, is_validation=False)
```



Examining the model

```
[22]: network.load_state_dict(torch.load('model.pt'))
```

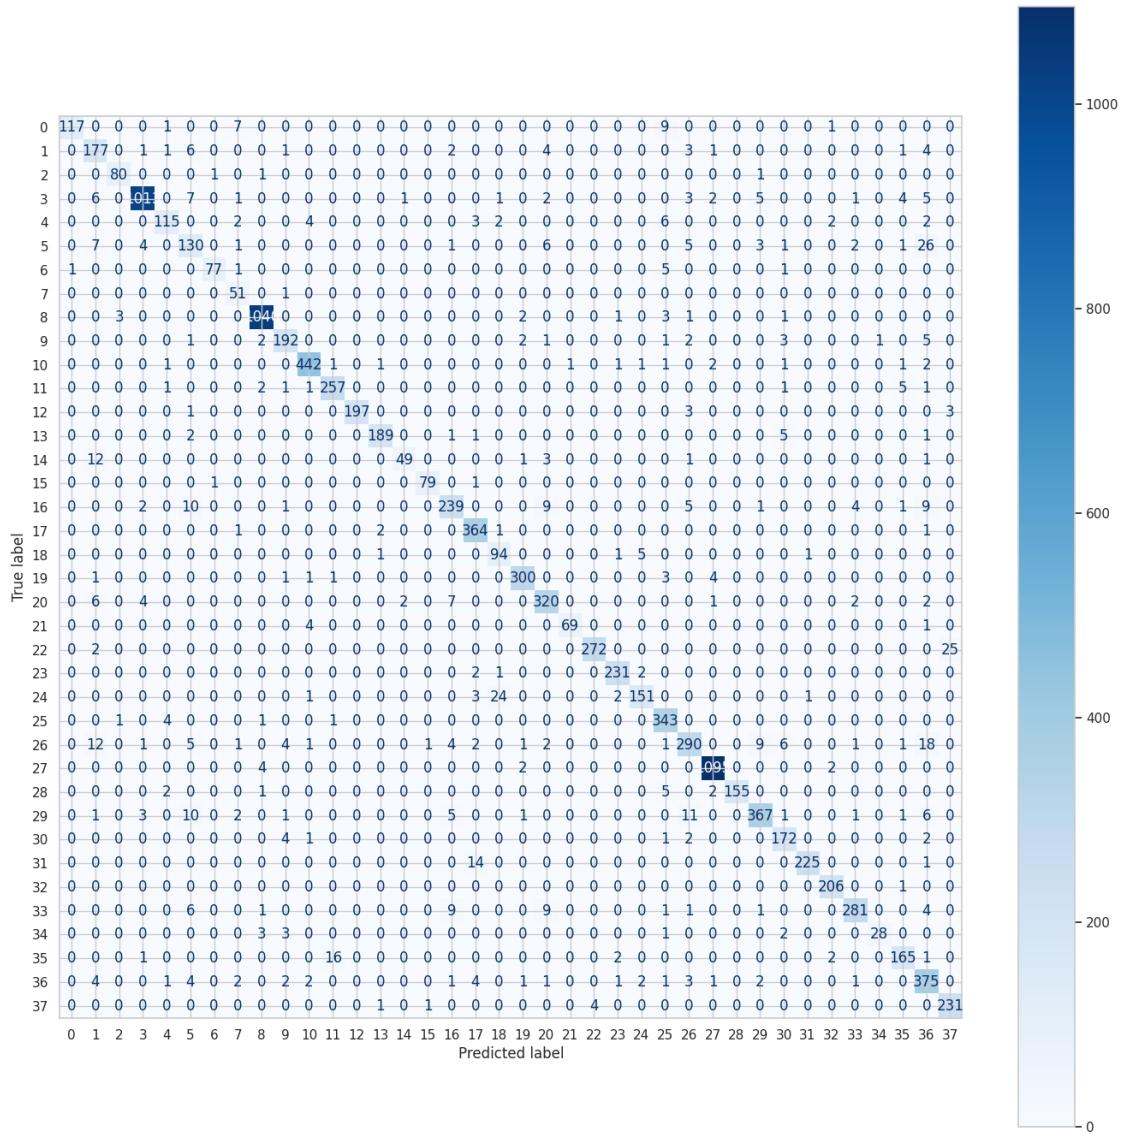
```
[22]: <All keys matched successfully>
```

```
[28]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
images, labels, probs, corrects = get_predictions(network, testloader, device)
pred_labels = torch.argmax(probs, 1)

print(f"There are {len(corrects)} correct predictions.")
```

There are 10863 correct predictions.

```
[29]: plot_confusion_matrix(labels, pred_labels)
```



```
[30]: # Now, extract and sort incorrect examples
incorrect_examples = []
for image, label, prob, correct in zip(images, labels, probs, corrects):
    if not correct:
        incorrect_examples.append((image, label, prob))

incorrect_examples.sort(reverse=True, key=lambda x: torch.max(x[2], dim=0).
                           values)
print(f"There are {len(incorrect_examples)} incorrect predictions.")
```

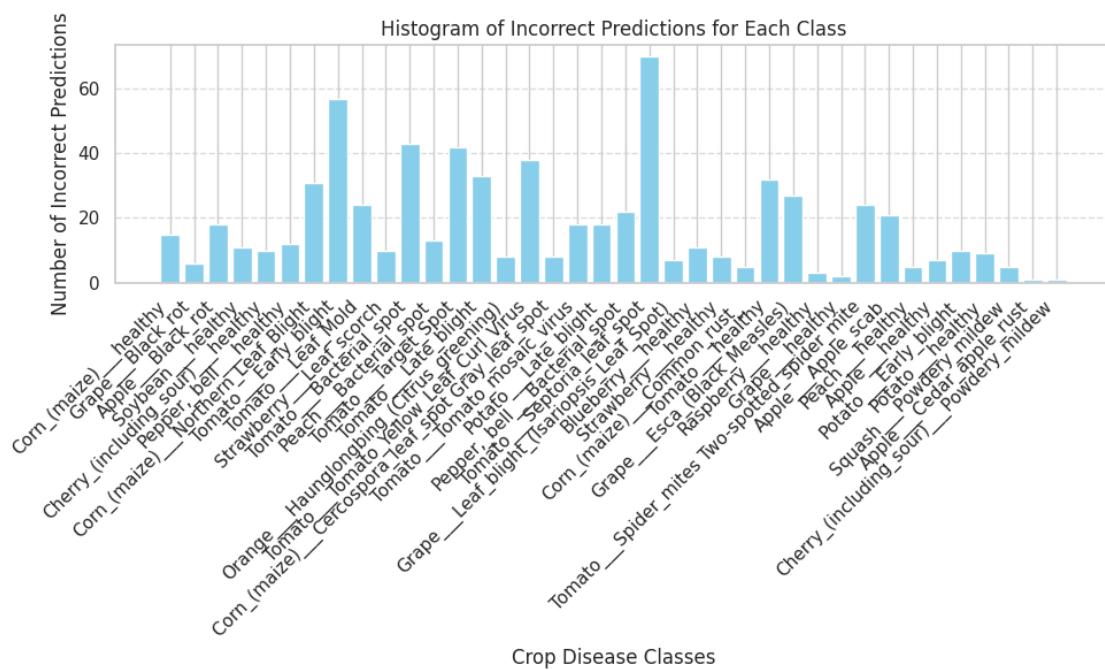
There are 685 incorrect predictions.

```
[31]: incorrect_counts = {}

# Count the occurrences of each class label in the incorrect examples
for _, label, _ in incorrect_examples:
    if label.item() not in incorrect_counts:
        incorrect_counts[label.item()] = 1
    else:
        incorrect_counts[label.item()] += 1

class_names = [train_set.classes[label] for label in incorrect_counts.keys()]

# Plotting the histogram
plt.figure(figsize=(10, 6))
plt.bar(class_names, incorrect_counts.values(), color='skyblue')
plt.xlabel('Crop Disease Classes')
plt.ylabel('Number of Incorrect Predictions')
plt.title('Histogram of Incorrect Predictions for Each Class')
plt.xticks(rotation=45, ha='right')
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()
```



```
[32]: data = {'Class Name': class_names, 'Incorrect Counts': list(incorrect_counts.  
    ↴values())}  
df1 = pd.DataFrame(data)
```

```
df1.head(5)
```

```
[32]:
```

	Class Name	Incorrect Counts
0	Corn_(maize)___healthy	15
1	Grape___Black_rot	6
2	Apple___Black_rot	18
3	Soybean___healthy	11
4	Cherry_(including_sour)___healthy	10

```
[33]:
```

```
class_counts = defaultdict(int)
for images, labels in testloader:
    for label in labels:
        class_counts[label.item()] += 1

class_labels = sorted(class_counts.keys())
counts = [class_counts[label] for label in class_labels]
colors = sns.color_palette("husl", len(class_labels))

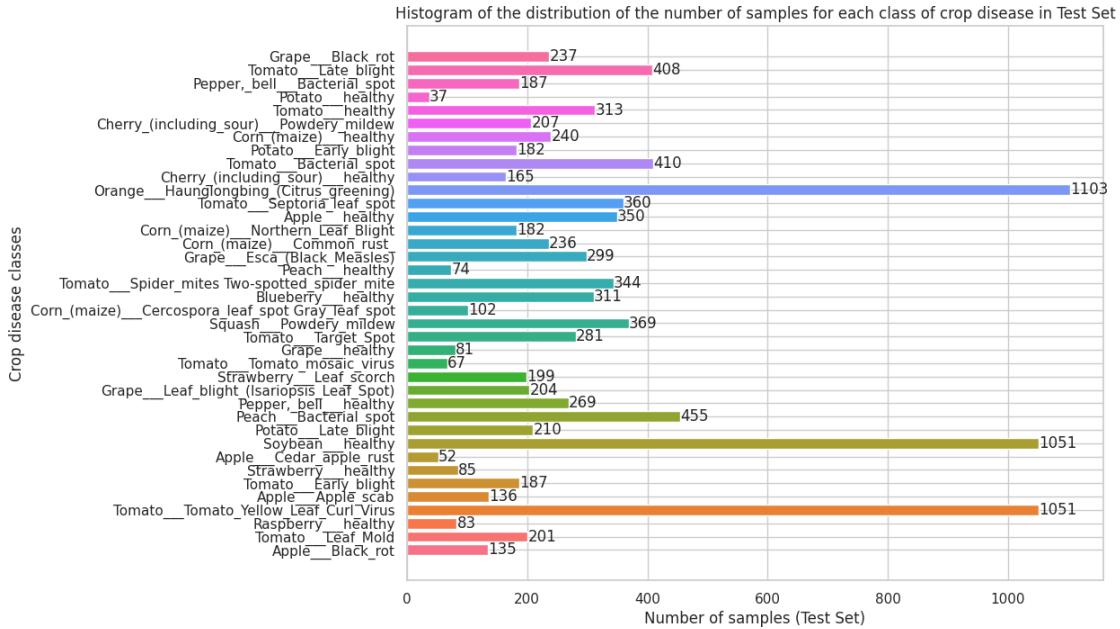
class_names = [train_set.classes[label] for label in class_labels]

sns.set(style="whitegrid")
plt.figure(figsize=(10, 8))
plt.barh(class_names, counts, color=colors)
plt.xlabel('Number of samples (Test Set)')
plt.ylabel('Crop disease classes')
plt.title('Histogram of the distribution of the number of samples for each class of crop disease in Test Set')

for i, (count, name) in enumerate(zip(counts, class_names)):
    plt.text(count, i, str(count), ha='left', va='center')

plt.show()
```

```
/usr/lib/python3.10/multiprocessing/popen_fork.py:66: RuntimeWarning: os.fork()
was called. os.fork() is incompatible with multithreaded code, and JAX is
multithreaded, so this will likely lead to a deadlock.
    self.pid = os.fork()
/usr/lib/python3.10/multiprocessing/popen_fork.py:66: RuntimeWarning: os.fork()
was called. os.fork() is incompatible with multithreaded code, and JAX is
multithreaded, so this will likely lead to a deadlock.
    self.pid = os.fork()
```



```
[34]: data = {'Class Name': class_names, 'Counts': counts}
df2 = pd.DataFrame(data)
df2.head(5)
```

```
[34]:          Class Name  Counts
0           Apple__Black_rot    135
1           Tomato__Leaf_Mold   201
2           Raspberry__healthy   83
3 Tomato__Tomato_Yellow_Leaf_Curl_Virus  1051
4           Apple__Apple_scab   136
```

```
[35]: merged_df = pd.merge(df1, df2, on='Class Name', suffixes=('_incorrect', '_test'))
merged_df.head(3)
```

```
[35]:          Class Name  Incorrect Counts  Counts
0  Corn_(maize)__healthy            15     240
1      Grape__Black_rot             6     237
2      Apple__Black_rot            18     135
```

```
[36]: # Calculate the rate of success for each class in percentage
merged_df['Success Rate (%)'] = ((merged_df['Counts'] - merged_df['Incorrect_Counts']) / merged_df['Counts']) * 100
merged_df
```

```
[36]:          Class Name  Incorrect Counts \
0  Corn_(maize)__healthy            15
```

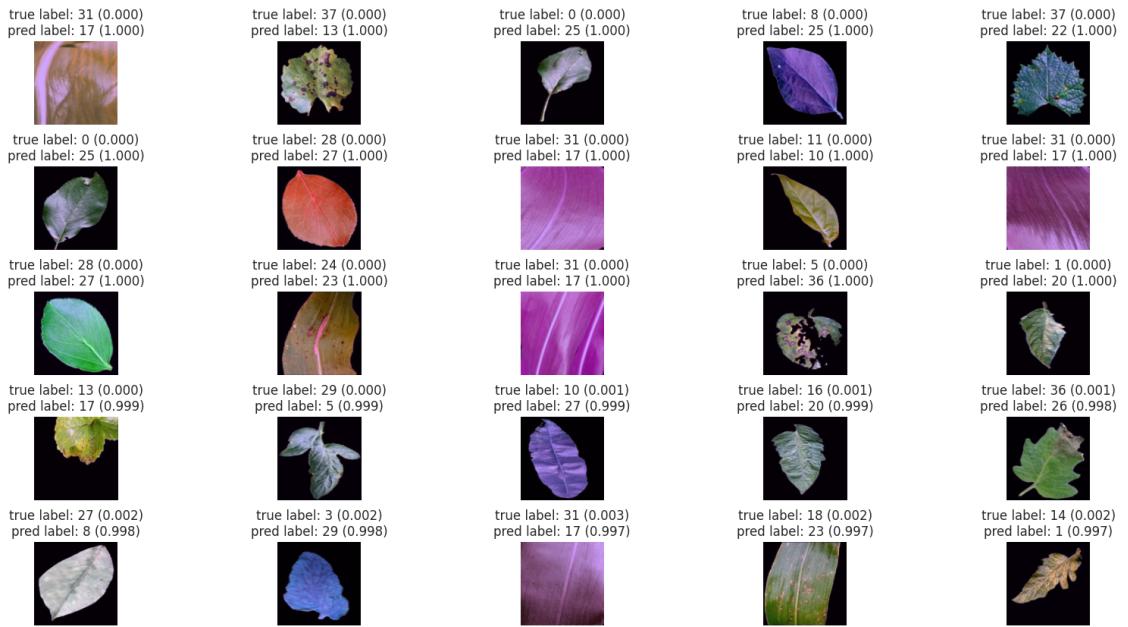
1	Grape___Black_rot	6
2	Apple___Black_rot	18
3	Soybean___healthy	11
4	Cherry_(including_sour)___healthy	10
5	Pepper,_bell___healthy	12
6	Corn_(maize)___Northern_Leaf_Blight	31
7	Tomato___Early_blight	57
8	Tomato___Leaf_Mold	24
9	Strawberry___Leaf_scorch	10
10	Tomato___Bacterial_spot	43
11	Peach___Bacterial_spot	13
12	Tomato___Target_Spot	42
13	Tomato___Late_blight	33
14	Orange___Haunglongbing_(Citrus_greening)	8
15	Tomato___Tomato_Yellow_Leaf_Curl_Virus	38
16	Corn_(maize)___Cercospora_leaf_spot_Gray_leaf_...	8
17	Tomato___Tomato_mosaic_virus	18
18	Potato___Late_blight	18
19	Pepper,_bell___Bacterial_spot	22
20	Tomato___Septoria_leaf_spot	70
21	Grape___Leaf_blight_(Isariopsis_Leaf_Spot)	7
22	Blueberry___healthy	11
23	Strawberry___healthy	8
24	Corn_(maize)___Common_rust_	5
25	Tomato___healthy	32
26	Grape___Esca_(Black_Measles)	27
27	Raspberry___healthy	3
28	Grape___healthy	2
29	Tomato___Spider_mites Two-spotted_spider_mite	24
30	Apple___Apple_scab	21
31	Peach___healthy	5
32	Apple___healthy	7
33	Potato___Early_blight	10
34	Potato___healthy	9
35	Squash___Powdery_mildew	5
36	Apple___Cedar_apple_rust	1
37	Cherry_(including_sour)___Powdery_mildew	1

Counts	Success	Rate (%)
0	240	93.750000
1	237	97.468354
2	135	86.666667
3	1051	98.953378
4	165	93.939394
5	269	95.539033
6	182	82.967033
7	187	69.518717

8	201	88.059701
9	199	94.974874
10	410	89.512195
11	455	97.142857
12	281	85.053381
13	408	91.911765
14	1103	99.274705
15	1051	96.384396
16	102	92.156863
17	67	73.134328
18	210	91.428571
19	187	88.235294
20	360	80.555556
21	204	96.568627
22	311	96.463023
23	85	90.588235
24	236	97.881356
25	313	89.776358
26	299	90.969900
27	83	96.385542
28	81	97.530864
29	344	93.023256
30	136	84.558824
31	74	93.243243
32	350	98.000000
33	182	94.505495
34	37	75.675676
35	369	98.644986
36	52	98.076923
37	207	99.516908

```
[39]: N_IMAGES = 25
```

```
plot_most_incorrect(incorrect_examples, N_IMAGES)
```



[]: